School of Computer Science and Information Technology
Lucerne University of Applied Sciences and Arts (Switzerland)

Demographic Biases in Dermatology Models
TODO: subtitle

Bachelor Thesis

presented to School of Computer Science and Information Technology of Lucerne
University of Applied Sciences and Arts (Switzerland) in consideration for the award of
the academic grade of *Bachelor in Computer Science.*

by

Nadja Stadelmann

from

Lucerne (Switzerland)

Declaration

Bachelor Thesis at Lucerne University of Applied Sciences and Arts
School of Computer Science and Information Technology

Title of Bachelor Thesis:	Demographic Biases inDermatology Models
Name of Student:	Nadja Stadelmann
Degree Program:	Bachelor in Computer Science
Year of Graduation:	2025
Main Advisor:	Dr. Ludovic Amruthalingam
External Expert:	Dr. Jürg Schelldorfer
Industry partner/provider:	Applied AI Research Lab

Code/Thesis Classification

☒ Public (Standard)
☐ Private

Declaration

I hereby declare that I have completed this thesis alone and without any unauthorized or external help. I further declare that all the sources, references, literature and any other associated resources have been correctly and appropriately cited and referenced. The confidentiality of the project provider (industry partner) as well as the intellectual property rights of the Lucerne University of Applied Sciences and Arts have been fully and entirely respected in completion of this thesis.

Place/Date, Signature _____

Submission of the Thesis to the Portfolio Database

Confirmation by the student

I hereby confirm that this bachelor thesis has been correctly uploaded to the portfolio database in line with the code of practice of the University. I rescind all responsibility and authorization after upload so that no changes or amendments to the document may be undertaken.

Place/Date, Signature _____

Expression of thanks and gratitude

Thanks to my family, relatives and friends for all the support given to finish this thesis. TODO: add thanks and gratitude Ludovic Amruthalingam Simone Lionetti - deputy Ludovic Pascal Baumann - LaTeX Philippe Gottfrois - information and work on PASSION project Proofreaders TODO: do you want to be mentioned with name or not?

Nadja Stadelmann, 2025

Summary

TODO: Your abstract here. The content of your thesis in brief.

Contents

Todo list

Alle Fakten (fundiertes Wissen Dritter) sind korrekt zitiert. Es werden verschiedene Zitierweisen verwendet und teilweise mehrere Interpretationen gegenübergestellt. Der gemeinsam definierte Zitierstil im Text, in Abbildungen und Tabellen sowie im Literaturverzeichnis wird korrekt und durchgängig angewendet. Eigene Leistungen (sowie Bewertungen) und Fremdquellen sowie Recherchen sind klar unterscheidbar.

Die erstellten Artefakte sind von sehr hoher Qualität. Das trifft u.a. auf Diagramme, Skizzen sowie Notationen (z.B. BPMN/UML) zu. Darstellungen sind einwandfrei, alle statistisch notwendigen Qualitätskriterien sind erfüllt. Beschriftungen etc. sind vorhanden, keine Einwände, Text und Bild stimmen beschreibend gut überein. Es wurden angemessene Dokumentationsmethoden und -arten korrekt verwendet. Vereinbarte Interview Transkripte, Beobachtungsprotokolle bzw. Zusammen-fassungen sind vorhanden. Daten, Ort, Kontext, Beschreibung, Zeilennummer, Verweise, Strukturen sind erkennbar, gut formatiert und korrekt mit dem Text/ der Analyse verknüpft. Alle Elemente und Themen sind im methodischen Teil/Text erklärt und verständlich, keine technischen oder strukturellen Einwände. Auch Zwischenanalysen, Zwischenschritte oder Gesamtauswertungen wurden durchgeführt, die Herkunft der Daten ist erkennbar und professionell aufbereitet.

Der Schreibstil aller Dokumente entspricht hohen Standards und enthält keine Übertreibungen oder unbegründete Beurteilungen. Die Sprache ist aussagekräftig, prägnant und präzise. Die Fachterminologie ist konsistent, d.h. für gleiche Gegenstände und Themen werden immer die gleichen Begriffe verwendet. Der Sprachgebrauch ist durchgängig geschlechtergerecht, einheitlich und sachlich.

List of Figures

List of Tables

Glossary

Fitzpatrick skin type A skin classifier based on the skins' reaction to ultraviolet light, developed by dermatologist Dr. Thomas Fitzpatrick (Gottfrois et al., 2024). ix, 4,

Jupyter Notebook Executable files, often used in ML to write Python code and add explanations in text form. 6

pediatric A medical term for infants, children and adolescents. 1, 4

proxy variable "one or more variables that encode the protected attribute with a substantial degree of accuracy" according to https://medium.com/bcggamma/practice-ai-responsibly-with-proxy-variable-detection-42c2156ad986. 5, 17

teledermatology dermatological care from a distance, supported by modern technology (Pala et al., 2020). 1, 4, 24

Acronyms

AI Artificial Intelligence. 3, 5, 8, 10, 21, 31, 36, 38, 40, 50, 51, 55, 56, 57, 58, 66

FST Fitzpatrick skin type. *Glossary:* Fitzpatrick skin type, 4, 5, 6, 13, 58, 60, 62

ML Machine Learning. vi, vii, 3, 4, 6, 8, 9, 11, 12, 15, 21, 32, 33, 34, 35, 37, 44, 45, 46, 47, 48, 49, 58

TODO: fix citations in glossary

1 Problem Statement

Welche Ziele, Fragestellungen werden mit dem Projekt verfolgt? Die Bedeutung, Auswirkung und Relevanz dieses Projektes für die unterschiedlichen Beteiligten soll aufgeführt werden. Typischerweise wird hier ein Verweis auf die Aufgabenstellung im Anhang gemacht.

In Sub-Saharan Africa dermatology treatment is inaccessible according to Gottfrois et al. (2024). There are fewer than one dermatologist available per one million people. Despite this, up to 80% of the children and adolescents in the area are affected by skin conditions. AI-based teledermatology promises to close this gap of specialists per case, for example by serving as a triage option. Potential patients could upload pictures to diagnostic dermatology AIs which can indicate whether the person should indeed visit a dermatologist or promote other treatment options. However, current dermatology AIs tend to fail to deliver accurate results for patients with highly pigmented skin tones. This is mainly due to demographic biases in existing AI models. The models are trained on established datasets which mainly feature low-pigmented skin. Therefore, the datasets lack representation of highly-pigmented skin, leading to AI models which do not generalize to the population in Sub-Saharan Africa (Gottfrois et al., 2024).

These biases result in unequal access to treatment and especially affect underrepresented groups. Such biased results must be avoided, especially in AI models which impact life-changing decisions (Mehrabi et al., 2021).

According to Diaz et al. (2022), demographic biases are especially important in dermatology. Demographic differences in patients influence the appearance of dermatological conditions. The differences in appearance can be developed depending on genetic factors, such as skin tone, age and sex (Diaz et al., 2022). Research showed, that in patients with lower socio-economic status the disease progression is more advanced at time of diagnosis, which in turn can lead to different appearances for the same disease (British Association of Dermatologists (BAD), 2021). Since the AI models use pictures as the inputs and can only learn to diagnose diseases according to their appearances in the data, the factors which affects the disease appearances must be considered when creating an inclusive dataset.

In order to overcome these issues, the PASSION research team founded the PASSION project. The projects vision is to make dermatology treatment accessible in Africa by enabling the AI-supported teledermatology for triage by reducing the demographic biases in the dermatology AI models. For this bias mitigation, the researcher collected a dataset in Sub-Saharan Africa, focusing on patients with highly pigmented skin and the most common regional pediatric skin condi-

tions. The PASSION dataset is complementary to existing datasets and improves their diversity. Further, the PASSION team trained an ResNet-50 model with the dataset which is refered to as PASSION model in this thesis. This model should serve as a benchmark model to assess other dermatology models in regards of fairness (Gottfrois et al., 2024). TODO: check sources, maybe, for the last sentence, the midterm protocol must be cited instead

So that the PASSION model can become an unbiased benchmark model, potential demographic biases in it must be reduced as far as possible. To reach this goal, demographic biases in the model as well as the limitation of the gathered dataset must be identified and mitigated. This thesis supports the PASSION team in this process. The main objective of the thesis is to assess the effectiveness of mitigation strategies to reduce demographic biases in context of PASSION.

- With the advent of telemedicine, developing countries are learning newer ways of leveraging their information and communication technologies (ICTs) to play an increasingly vital role in the health care industry. Telemedicine is defined as a health care delivery mechanism where physicians and other medical personnel can examine patients remotely using information and telecommunication technologies (ICTs; Bashshur, Sanders, and Shannon, 1997). (**Kifle_2024**)

- AI systems can be used in many sensitive environments to make important and life-changing decisions; thus, it is crucial to ensure that these decisions do not reflect discriminatory behavior toward certain groups or populations (Mehrabi et al., 2021).

- There are clear benefits to algorithmic decision-making; unlike people, machines do not become tired or bored [45, 119], and can take into account orders of magnitude more factors than people can. However, like people, algorithms are vulnerable to biases that render their decisions "unfair" [6, 121]. In the context of decision-making, fairness is *the absence of any prejudice or favoritism toward an individual or group based on their inherent or acquired characteristics*. Thus, an unfair algorithm is one whose decisions are skewed toward a particular group of people. (Mehrabi et al., 2021).

- it is important for researchers and engineers to be concerned about the downstream applications and their potential harmful effects when modeling an algorithm or a system (Mehrabi et al., 2021).

- We should think responsibly, and recognize that the application of these tools, and their subsequent decisions affect peoples' lives; therefore, considering fairness constraints is a crucial task while designing and engineering these types of sensitive tools (Mehrabi et al., 2021).

2 State of Research

Bezogen auf die eigenen Zielsetzungen und Fragestellungen soll aufgezeigt werden, wie andere dieses oder ähnliche Probleme gelöst haben. Worauf können Sie aufbauen, was müssen Sie neu angehen? Wodurch unterscheidet sich Ihre Lösung von anderen Lösungen? Für wissenschaftlich orientierte Arbeiten sei hier explizit auf (Balzert, S. 66 ff) verwiesen. Relevante, aktuelle und fundierte Fachliteratur wurde identifiziert, kritisch geprüft und verwendet. Die Begriffe der Fragestellung sind definiert und referenziert. Der gesamte Kontext ist verknüpft und eine Abgrenzung wurde vorgenommen. All dies ist in einer leicht verständlichen Struktur formuliert und überprüft.

In this chapter a review over the existing work in the field of bias mitigation in Artificial Intelligence (AI). The main focus lies on a literature review of existing papers from other researchers in this area, highlighting the key findings which are connected to this thesis. Bias mitigation in AI has already been investigated by different researchers, who crafted fitting mitigation methods TODO: citation?. This thesis aims to assess those existing methods in the context of PASSION.

Therefore, this chapter first presents an overview over the PASSION project based on the PASSION paper and dataset. Then, the general knowledge in the literature about existing biases, fairness metrics and mitigation methods is summarized. The review process was divided in two main contexts: ML in general and ML in dermatology. This approach ensures that the technical and dermatological perspectives are considered when applying the knowledge to PASSION. The tables in this chapter indicate which points where found in which context. This is important, since what may be an issue in general might not be relevant for a specific use case or vice versa. For example, in theory, all age groups should be represented in datasets to account for demographic diversity. However, for car insurance, the age representation is not important, because age does not affect how well a driver can drive TODO: either cite this example from the expert or find another example related to dermatology.

The various studies present different bias sources and suggest diverse methods to mitigate them. During the literature review, several biases and mitigation methods were identified that may be relevant to the PASSION project. Since it is not feasible to assess all of them during the duration of this thesis, the thesis focuses on those which are related to skin type, age and gender. The chosen methods are explained in chapter 4 Methods. The other items are passed to the PASSION research team as a list for further investigation. The list can be found in the appendix TODO: add link.

2.1 PASSION for Dermatology

This section provides an overview of to the PASSION project regarding its medical scope and technical components.

While the overall goal remains to improve the accessibility of dermatological care by building fair and inclusive AI systems, PASSION specifically addresses common pediatric skin conditions in Sub-Saharan Africa. To create a dataset which represents patients with highly pigmented skin, they collected data from patients with Fitzpatrick skin types (FSTs) III to VI. Based on this dataset, the PASSION team fine-tuned a ResNet-50 model using transfer learning. With the dataset and trained model, the researchers published data analysis scripts and initial insights on the model performance in a MICCAI TODO: add to glossary publication (Gottfrois et al., 2024).

For the purpose of this thesis, it is essential to understand the dataset's metadata, the architecture and fine-tuning process of the PASSION model and which bias mitigation methods have already been applied. The dataset can influence which biases could arise in the model or rather which ones can be measured. The labels which should be predicted and the model architecture give insight into the ML task. All this information affect which mitigation methods are feasible to be used for the project. TODO: add sources

2.1.1 PASSION Dataset

The PASSION dataset contains data from patients from four African countries in dermatology clinics. It contains 4901 images of 1653 dermatology cases with the corresponding demographic and clinical information, see Table 2.1. There is one record per patient and one or more corresponding images which are linked to the record through the filename. The images are taken with mobile phones, in order to train the models on a teledermatology-like setup regarding image quality (Gottfrois et al., 2024).

The datasplit is a predefined 80/20 stratified development-test split at patient level. This ensures reproducibility, fair comparison, while preventing information leaking (Gottfrois et al., 2024). Stratified splitting is a method to split imbalanced datasets without changing the original class distribution within the subset. This is important to maintain the representativeness of the data, especially for minority classes (Baldé, 2023). TODO: add better source for the stratified datasplit (and the other stuff where medium was used).

Access to the dataset can be requested via https://passionderm.github.io/ (Gottfrois et al., 2024).

Metadata Attribute	Data Type	Description
subject_id	string	Participant's unique identifier
country	string	Country of data origin
age	integer	Age of the participant in years
sex	m/f/o	Gender of the participant
fitzpatrick	integer	FST
body_loc	string (list; null-able, semicolon-separated)	Specific affected body locations
impetig	0/1	Presence of impetigo (1=present), may occur alone or with other conditions, affects the treatment options for coexisting conditions
conditions_PASSION	Eczema, Scabies, Fungal, Others	Primary diagnosed skin condition

Table 2.1: PASSION dataset - metadata attributes and descriptions (Gottfrois et al., 2024)

For this thesis, the attributes *fitzpatrick*, *sex* and *age* are relevant, as they cover relevant demographic information from the patients which can be used to identify demographic biases. The skin type is directly relevant because it can affect disease presentation and therefore the model's learning process to detect the condition TODO: add citation, also from mail. According to Philippe Gottfrois, the main author of the PASSION paper, the same condition looks similar regardless of a patient's age and sex. However, the prevalence of conditions varies based on these factors TODO: add citation, also from mail. Therefore, the data distribution of age and sex could become relevant for an inclusive AI model, because these demographic factors influence the likelihood of a condition being represented in the dataset. In this thesis, the attributes will be used to identify potential biases in the PASSION model. Further, it will be evaluated whether the data distributions for these attributes indeed are relevant for any biases TODO: check in the end if I really did that.

The attribute *country* seems to be another demographic factor. Since it is the country where the data collection took place TODO: cite mail, the only demographic information it could indicate is the geographic location where the patient was diagnosed. This could potentially be used as proxy variable for where the patient lives, which could be relevant for disease prevalance TODO: cite this. However, this could lead to false conclusions which is why this label will not be used in this thesis. The PASSION team should investigate the need for this label and possible enhancement. Also, they should state the meaning of the label clearer in the dataset description. TODO: add this to PASSION investigation list

The labels *impetig* and *conditions_PASSION* are the ones which the PASSION

model learns to predict. They can appear independently of each other. Therefore, the ML task for PASSION it is a multiclassification task.

2.1.2 PASSION Data Analysis Scripts

The PASSION team provides a Jupyter Notebook with code examples and analysis scripts. They are listed in Table 2.2 together with their relevance to this thesis. The most relevant scripts are those related to demographic distributions of the chosen attributes, since they help identifying potential data imbalances. Scripts that lay the foundation for further analysis are somewhat relevant, while all other scripts are irrelevant for this thesis.

TODO: add the distribution analysis from the presentation
TODO: move tbl to appendix

Script Title	Description	Relevance - Reasoning
Distribution of FSTs	Counts and visualizes the skin type distribution	**High** - Insight into demographic distributions
Regrouping Malawi and Tanzania to EAS	Data aggregation due to dataset size and geographical proximity	**Medium** - Might impact interpretation of the results of the following scripts
Linking CSV Data with Image Files	Mapping between data records and images.	**Medium** - Basis for other analyses
Extracting and Comparing Subject IDs	Dataset verification regarding completeness	**Low** - No insight in regards of demographic distribution
Conditions by Country	Correlation between clinical conditions and country	**Low** - The attribute *country* is out of scope of this thesis
Body Localizations by Conditions	Correlation between the condition and primarily affected body parts	**Low** - No insight in regards of demographic distribution
Impetigo Cases	Total count of impetigo cases and proportion to all cases	**Low** - No insight in regards of demographic distribution[*]

[*] Research is divided which demographic factors play into the prevalence of impetigo (Romani et al., 2017; Aleid et al., 2024).

Table 2.2: PASSION dataset - existing analysis scripts (Gottfrois et al., 2024)
TODO: decide on a table style

2.1.3 PASSION Model

The model architecture is a ResNet-50 model which is pretrained on ImageNet. The model was fine-tuned by replacing the last fully connected classification layer

with a dropout layer with a 0.3 dropout rate followed by batch normalization. The class activation is done by a single linear layer. To minimize the weighted cross-entropy loss, Adam optimization is used. For improved generalization and to avoid overfitting, data augmentations were applied. The methods used were random resizing, cropping, flipping, and rotating Gottfrois et al. (2024). TODO: add individual citations for ResNet-50, ImageNet, Adam optimization, weighted cross-entropy loss.

TODO: add information regarding how many folds are there, how is the data split, ...

2.1.4 PASSION Experiments

The PASSION team conducted various experiments to evaluate the classifiers on the test set with the following schemes (Gottfrois et al., 2024):

- Performance for skin condition prediction

- Performance for impetigo detection

- Generalization from two centers to a wider population (test set contains data from the known centers and one unknown center)

- Generalization from different age groups (test set contains data from the known age groups and one unknown)

- Subject level analysis over the predictions of multiple pictures, using majority voting

The code for those experiments is available in the PASSION evaluation GitHub repo. These repo can serve as a starting point, since reproducing the results helps to verify that the provided setup works the same on my side. Also, they can be used as examples for further experiments. TODO: mention which ones I really used why for the thesis and move the others to the appendix

The paper indicates lower performance when evaluating the model on a subject level (performance per case/patient) rather than a sample level (performance per image). The authors emphasize the importance of assessing classifier performance on both levels for completeness (Gottfrois et al., 2024). Therefore, the subject level performance should also be considered during this thesis.

2.1.5 Limitations

TODO: maybe move to execution phase TODO: write in more details - multiple executions showed inconsistent results for the different group evaluations on the same model checkpoint. It turned out that the metadata linkage did not work consistently. The issue was resolved now by providing the img name in the dataloader and link the metadata directly from the source file instead of using the indexes. prob related to different shufflings between dataloader and metadataloader

2.2   Bias

This chapter provides an overview of biases and related demographic characteristics mentioned in ML- and dermatology-related research. It also explains their relevance for PASSION.

Algorithmic decisions made by AI systems can directly affect peoples' lives. In healthcare applications such as PASSION, these decisions are especially sensitive, as they influence diagnoses and treatment outcomes. Diverse studies have shown that AI application's decisions can hold biases that affect underrepresented groups. This leads to unfair or even harmful consequences. Therefore, it is essential for AI engineers to identify, address and mitigate such biases in order to develop fair applications. This requires to understand what bias is in general, what concrete biases exist, and where they originate (Mehrabi et al., 2021).

2.2.1   Definition of Bias in ML

In the context of ML, bias can be defined as *a systematic error that causes a model or estimator to consistently deviate from the true value or relationship* (Delgado-Rodríguez & Llorca, 2004; Taylor, 2023). In practice, this often results in models that make less accurate predictions for specific subgroups within the population TODO: cite this.

TODO: ChatGPT Text starts here, make sure it is cited correctly

2.2.2   Demographic Biases in the Context of Dermatology

Biases in dermatology in general can lead to unequal outcomes for different groups, which can result in unfair outcomes for certain groups. Demographic biases are particularly relevant in the context of dermatology AIs, as they can cause differences in diagnostic accurracy and treatment outcomes among different demographic (sub-)groups. From the literature review, three main ways have been identified in which demographic differences may introduce bias in dermatology ML models: TODO: cite all that, from presentation

- **Disease Presentation**. *Skin type* affects how diseases appear on the skins. As Gottfrois notes, "any condition linked to inflammation is less visible if the skin is more pigmented" TODO: cite mail from philippe. This directly influences training and evaluating image-based ML models like those used in PASSION. For example, a model trained predominantely on images with low pigmented skin may perform poorly on images of highly pigmented skin.

- **Disease Prevalence**. Factors such as *age* and *sex* do not tend to affect disease presentation, but they can influence disease prevalence TODO: cite mail from philippe. Therefore, these factors could introduce bias if certain conditions are underrepresented in the dataset due to demographic imbalances. TODO: consider adding smt like the car driver example here, indicating that it is not necessarily a problem due to the same disease presentation

- **Access to Healthcare**. *Socioeconomic status* or *geographic location* can also introduce bias. Research shows that patients with lower socioeconomic status are often diagnosed at later stages of the disease, which may alter the visual presentation of the disease. If such cases are missing in training data, the model will fail to recognise them, leading to misdiagnosis. TODO: add example for geographic location?.

To build a robust and fair ML model, it is essential to identify and address this sources of biases (**Mehrabi2022**).
TODO: if citing is an issue: check the comment

2.2.3 Types of Biases and Their Relevance for PASSION

The literature describes numerous types of bias. Over 60 were identified during this research. These were grouped into categories to provide an overview. Their relevance for PASSION context was assessed.

Among them, **sampling biases** and **representation biases** are particularly relevant, as they relate directly to the inclusion or exclusion of demographic subgroups in the dataset. For example, **ascertainment bias**, a subtype of sampling bias, occurs when parts of the target population are unintentionally excluded. A common example is healthcare studies conducted in public hospitals only, which excludes patients from higher socioeconomic backgrounds who visit private clinics. This skews the data and can lead to incorrect conclusions, such as overestimating disease prevalence in specific groups.

Other relevant categories include **medical biases** and **imaging biases**, especially in the teledermatology setting of PASSION. These include clinical labeling errors, variations in image quality or lighting conditions which lead to bias.

This thesis focuses on the most relevant bias types. An extensive list is provided in the appendix and will be shared with the PASSION team for further evaluation. TODO: add the biases here
TODO: use this template also in the appendix, ask ChatGPT to fill it out

2.2.4 Template for Describing Relevant Biases

To systematically describe the most relevant biases for the PASSION project, the following template is used:

- **Name:** Ascertainment Bias

- **Category:** Sampling Bias

- **Definition:** A systematic distortion arising from the method by which participants or data are selected for inclusion.

- **Example:** Studies on STD prevalence conducted only in public clinics may overlook patients from higher-income backgrounds who go to private practitioners.

- **PASSION Relevance:** If PASSION's dataset is composed mostly of patients from certain types of clinics, it may not generalise well to other socioeconomic groups.

- **Mitigation Strategy:** Ensure that data is collected from a diverse range of sources, including both public and private healthcare facilities.

This structure will be used to explain 5–10 selected biases in detail in the following sections. The complete list of identified biases and their classifications can be found in the appendix.

TODO: consider to use those chapters for fairness introduction

2.2.5 Bias Detection and Fairness Assessment

Detecting bias in datasets and model outputs is the first step toward mitigation. However, bias detection alone is not sufficient. It is also necessary to evaluate the fairness of the model's predictions. This should be done across different demographic groups to ensure equitably performance in real-world applications.

An initial analysis of the PASSION dataset indicates an overrepresentation of male children. This imbalance should be further investigated, both in terms of dataset composition and model performance.

2.2.6 Focus for this Thesis

TODO: check that there is no duplication between PASSION dataset feature description and here Due to time constraints, this thesis focuses on three protected characteristics: *skin type*, *age* and *sex*. These were selected based on their presence in the PASSION dataset and their influence on dermatological diagnosis and disease prevalence. Other potentially relevant features, such as *country of origin* and *socioeconomic status*, should be evaluated in future work by the PASSION team.

TODO: continue here for fairness and so on

According to the Cambridge English dictionary, bias can be defined as "the fact of preferring a particular subject or thing" or even as "the action of supporting or opposing a particular person or thing in an unfair way, because of allowing personal opinions to influence your judgment" (**Cambridge_2025**). While an AI does not hold personal opinions, the judgment of an AI algorithm is influenced by the underlying data and how the algorithm uses this data. The data and even the algorithm itself can hold biases, which affects the final outcome and potential lead to unfair decisions. In decision-making, fairness means "absence of any prejudice or favoritism toward an individual or group based on their inherent or acquired characteristics" - or in other terms, a fair algorithms decisions are not scewed toward a group of people (Mehrabi et al., 2021). TODO: Do I need to say more about this? If so, add infos from section Discrimination vs. Biases

- The Cambridge English dictionary defines bias as "the action of supporting or opposing a particular person or thing in an unfair way as a result of allowing personal opinions to influence your judgement."[1] However, statistical

bias is defined as any systematic error in the determination of the association between exposure and disease.2 (Chakraborty, 2024)

- These biased predictions stem from the hidden or neglected biases in data or algorithms (Mehrabi et al., 2021).

- There are clear benefits to algorithmic decision-making; unlike people, machines do not become tired or bored [45, 119], and can take into account orders of magnitude more factors than people can. However, like people, algorithms are vulnerable to biases that render their decisions "unfair" [6, 121]. In the context of decision-making, fairness is *the absence of any prejudice or favoritism toward an individual or group based on their inherent or acquired characteristics*. Thus, an unfair algorithm is one whose decisions are skewed toward a particular group of people. (Mehrabi et al., 2021).

2.2.7 Bias Sources

The general ML lifecycle consists of data gathering, training the algorithm and the user interaction with the trained model. Now, while data gathering, biases can arise either through the collection process or it is already inherited in the available data. Further, depending on the algorithm design, during training, the existing bias in the data can be amplified and new bias can be introduced. Lastly, the result of the algorithm can affect the user experience on inference which can lead to further bias amplification. This generates a feedback loop between the biases in each step of the ML lifecycle which can make it hard to identify the original bias source. The feedback loop is illustrated in Figure 2.1, which also shows first bias definitons, which were categorized according to this feedback loop (Mehrabi et al., 2021).
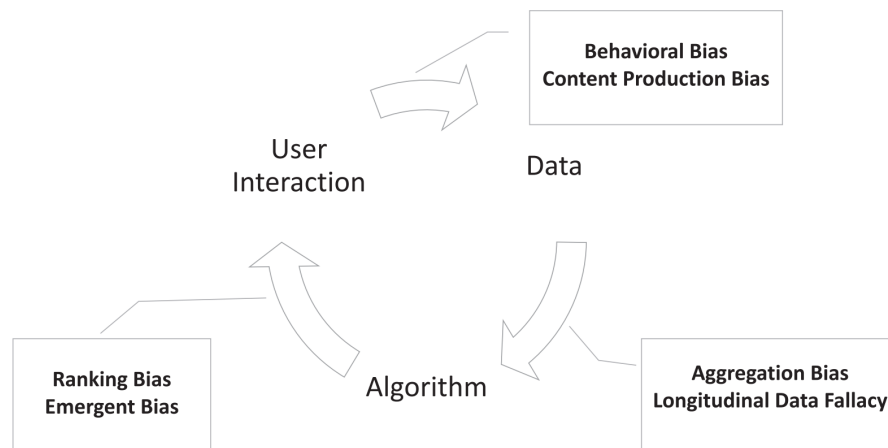

Figure 2.1: Bias definitions in a ML lifecycle (Mehrabi et al., 2021).

- two potential sources of unfairness in machine learning outcomes - those that arise from biases in the data and those that arise from the algorithms ...

we observe that biased algorithmic outcomes might impact user experience, thus generating a feedback loop between data, algorithms and users that can perpetuate and even amplify existing sources of bias (Mehrabi et al., 2021).

- The loop capturing this feedback between biases in data, algorithms, and user interaction is illustrated in Figure 1. We use this loop to categorize definitions of bias in the section below (Mehrabi et al., 2021)

2.2.8 Bias Types

The Table 2.3 aims to provide an overview over what kind of biases exist according to research. The more detailed categories listed in the table try to capture similar kind of biases. This thesis follows roughly the categorization of Mehrabi et al. (2021). Some biases might acctually fit in multiple categories. The definition of the categories including examples of specific biases follows.

TODO: check the c citations in the following chapters

Bias	Mentioned in Context of	
	ML	Dermatology
Data Biases		
Sampling Biases	X[1,2,3]	X[4]
Representation Biases	X[1]	X[5,6]
Measurement Biases	X[1,3]	X[4,6]
Research Biases	X[7]	X[4]
Feature Representation Biases	X[1,3]	X[4]
Imaging Biases		X[5]
Medical Biases	X[8]	X[4]
Temporal Data Biases	X[1]	X[4]
Algorithmic Biases		
User-Algorithm Interaction Biases	X[1]	
External Influence Biases	X[1]	X[4]
User Biases		
Cognitive Biases	X[1,7]	X[4]
Behavioral Biases	X[1,3]	X[4,5]
Publication Biases		X[4]
Medical Biases	X	X[4]

[1] (Mehrabi et al., 2021)   [4] (Chakraborty, 2024)   [7] (Mester, 2017)
[2] (HP, 2022)   [5] (Young et al., 2020)   [8] (Delgado-Rodríguez & Llorca,
[3] (Mester, 2022)   [6] (Montoya et al., 2025)   2004)

Table 2.3: Bias categories - grouped according the ML lifecycle of Mehrabi et al. (2021)

TODO: Feedback Astrid: zu viele Unterkapitel -> anders strukturieren

2.2.8.1 Data Biases

Sampling Biases

When gathering data, it's usually not possible to gather the data of a whole population. Instead, the data is gathered by sampling. A sample is a subgroup of individuals from the population. To get unbiased results, this sampling process should represent the true population, with a low sampling error (HP, 2022). This is often achieved with randomized samples. With non-random sampling processes, sampling bias arises. The consequence is, that the insights of one sampled population may not generalize with insights on another sampled popluation (Mehrabi et al., 2021).

Those biases can be introduced with a flawed sampling process:

- **Sampling bias**, due to nonrandom sampling of subgroups, leading to poor generalization (Mehrabi et al., 2021)

- **Selection bias**, working only on specific subset of the population which is not representative (Mester, 2022; Chakraborty, 2024)

- **Systematic selection bias**, chosen samples differ dramatically from the representative populations; e.g. in dermatology, when only the most severe patient data gets included (**c5**; **c6**; **c33**; Chakraborty, 2024)

- **Ascertainment bias**, tendency to exclude segments from the population due to e.g. cultural differences, such as which patient segment goes to government clinics vs. private clinics (usually influenced by socioeconomic status) (**c5**; Chakraborty, 2024)

- **Availability bias**, focus on widely available data instead of most representative data (**c9**; **c10**; Chakraborty, 2024**<empty citation>**)

- **Survivorship bias**, focus only on pre-selected data, ignoring the initial data-points which got filtered out (Mester, 2022).

Potential Biases in PASSION PASSION tries to reduce sampling bias in dermatology against high pigmented skin. PASSION might introduce (systematic) selection bias or Ascertainment bias, if in the dermatology centers only sickest / more severe patients are seen as indicated by Chakraborty (2024) PASSION inherits availability bias as it is using FST scale. Survivorship bias could be relevant for PASSION, if dermatology diseases could be lethal. Further, all patients which are not able to go to one of the dermatology centers which were used in PASSION could be considered to left out by survivorship bias.

used

- Sampling Bias. Sampling bias is similar to representation bias, and it arises due to nonrandom sampling of subgroups. As a consequence of sampling bias, the trends estimated for one population may not generalize to data collected from a new population. (Mehrabi et al., 2021). This is what the PASSION dataset tries to improve

- Selection bias - wrong sampling method, working on a specific subset of audience; usually by working only with data that is easy to access (Mester, 2022; Mester, 2017) - statistical bias

- Selection bias: Since it is not possible to work with large populations, for most dermatological studies, samples are chosen that are said to be representative of the original population. In selection bias, the selected subgroups are not representative of their original population. A variation of this is systematic selection bias, where samples chosen differ dramatically from their representative populations. Our experience suggests, such selection bias occurs more commonly in studies conducted in regional referral centers where only the sickest or more severe patients are usually seen. For example, a study compared the efficacy of thalidomide vs. prednisolone in hospitalised patients of erythema nodosum leprosum. It derived that thalidomide was more efficacious than steroids in erythema nodosum leprosum. Such findings cannot be generalised to all erythema nodosum leprosum since patients admitted to a regional referral center will likely have more severe disease.[5,6,33] (Chakraborty, 2024)

- Availability bias: More emphasis is placed on widely available data than scantily available data. A classic example is the use of antihistamines in pregnancy dermatoses, where nearly all standard books recommend first-generation antihistamine chlorpheniramine because more data is available.[9] 10. (Chakraborty, 2024) - dermatology

- Survivorship bias (Mester, 2022; Mester, 2017) - statistical bias

- Ascertainment Bias: This bias is commonly encountered in venereology practice. It is defined as a bias due to the tendency of some segments of the target population to get excluded due to cultural and other differences. For example, in most venereology clinics in government setups, studies show that venereal diseases are commoner in lower socioeconomic status. One reason might be that the higher socioeconomic status people tend to go to private practitioners and thereby get excluded from government-run clinics.[9,10] Allocation concealment and blinding are good ways to avoid this. 5. (Chakraborty, 2024) - healthcare

even more extensive

- Selection bias is again divided into two types endogenous selection bias and exogenous selection bias. The best example of endogenous selection bias in dermatology is the inclusion of non-response. If a trial tests the efficacy of a particular biologic in psoriasis, the response is usually collected from trial participants via postal services. Certain participants will not respond, although they might have substantially improved. Their exclusion will result in significant differences in efficacy evaluation.[33] Exogenous selection bias results when both treatment and outcome result from dependency on

an external variable that is not controlled. For example, if sunlight exposure is not controlled, it will influence both the intervention and control groups since psoriasis is a photosensitive (and photoexcerbated) dermatosis. (Chakraborty, 2024) - dermatology

- survivorship bias - World War II planes (**Silfwer__2017**) - https://doctorspin.org/media-psychology/psychology/survivorship-bias/

Representation Biases

TODO: still describe this category

Those biases can be introduced :

- **Representation bias**, non-representative sample lead to missing subgroups or other representation anomalies, which can be harmful to downstream applications. Popular ML datasets suffer from representation bias (Mehrabi et al., 2021; Shankar et al., 2017)

- **Population Bias**. Population bias arise when statistics, demographics and characteristics in the sample differ from the target population (Olteanu et al., 2019). The data it creates is non-representative for the target population (Mehrabi et al., 2021).

- **Aggregation bias** occurs, when "false conclusions are drawn about individuals from observing the entire population". It doesn't matter, whether the subgroups are represented equally in the training set, any generalized assumptions can result in aggregation bias (Mehrabi et al., 2021). In medicine, diseases can present themselves differently across genders and ethnicities (Suresh & Guttag, 2021). Therefore, diagnostic models need to incorporate those differences to mitigate aggregation bias (Mehrabi et al., 2021).

- **Simpson's Paradox** is a type of aggregation bias, which arises in heterogeneous data analysis. Observed associations disappear or reverses in the subgroup data (Mehrabi et al., 2021).

Potential Biases in PASSION PASSION tries to mitigate representation bias, by including more FST skin types - however, it could introduce other representation biases Aggregation bias and Simpson's Paradox could potentially be an issue when the analyzed skin diseases present themselves differently in patients based on their genetics

used

- Representation Bias. Representation bias arises from how we sample from a population during data collection process (Suresh & Guttag, 2021). Non-representative samples lack the diversity of the population, with missing subgroups and other anomalies (Mehrabi et al., 2021).

- Popular machine-learning datasets that serve as a base for most of the developed algorithms and tools can also be biased—which can be harmful to the downstream applications that are based on these datasets. ... In (Shankar et al., 2017), researchers showed that these datasets suffer from representation bias and advocate for the need to incorporate geographic diversity and inclusion while creating such datasets. (Mehrabi et al., 2021)

- Population Bias. Population bias arises when statistics, demographics, representatives, and user characteristics are different in the user population of the platform from the original target population (Olteanu et al., 2019). Population bias creates non-representative data. ... More such examples and statistics related to social media use among young adults according to gender, race, ethnicity, and parental educational background can be found in (Hargittai, 2007). (Mehrabi et al., 2021)

- Aggregation Bias. Aggregation bias (or ecological fallacy) arises when false conclusions are drawn about individuals from observing the entire population. An example of this type of bias can be seen in clinical aid tools. Consider diabetes patients who have apparent morbidity differences across ethnicities and genders. Specifically, HbA1c levels, that are widely used to diagnose and monitor diabetes, differ in complex ways across genders and ethnicities. Therefore, a model that ignores individual differences will likely not be well-suited for all ethnic and gender groups in the population (Suresh & Guttag, 2021). This is true even when they are represented equally in the training data. Any general assumptions about subgroups within the population can result in aggregation bias. (Mehrabi et al., 2021). –> could also be important for dermatology issues!!!

  - Simpson's Paradox. Simpson's paradox is a type of aggregation bias that arises in the analysis of heterogeneous data [18]. The paradox arises when an association observed in aggregated data disappears or reverses when the same data is disaggregated into its underlying subgroups (Fig. 2(a)). ... After analyzing graduate school admissions data, it seemed like there was bias toward women, a smaller fraction of whom were being admitted to graduate programs compared to their male counterparts. However, when admissions data was separated and analyzed over the departments, women applicants had equality and in some cases even a small advantage over men. The paradox happened as women tended to apply to departments with lower admission rates for both genders. Simpson's paradox has been observed in a variety of domains, including biology [37], psychology [81], astronomy [109], and computational social science [91].(Mehrabi et al., 2021).

Measurement Biases

How features are chosen, used and measured can lead to biases (Mehrabi et al., 2021; Suresh & Guttag, 2021).

Examples for such biases are:

- **Measurement bias** in general, e.g. using mismeasured proxy variables lead to misinterpretations of the outcome (Mehrabi et al., 2021)

- **Observer bias** is a subconscious bias which can occur in different forms. Either, researchers projects their own expectations on the research and influence the testers accordingly (Mester, 2022). In other cases, different observes report the same observation differently (**c29**; **c26**; Chakraborty, 2024)

- **Annotator bias** is a special form of observer bias. The labeling process of human annotators can be influenced by lots of factors (e.g. personal background, social context) and even minor design choices (e.g. scale order, image context). This can introduce inconsistencies when labeling the data (Montoya et al., 2025)

- **Recall bias**. This bias occurs when queried individuals do not remember things correctly, due to humans selective memory. This can cause misinterpretation, for example when analyzing causes and effects of behaviour on certain diseases in medicine (Mester, 2022**c3-6**; **c2**; Chakraborty, 2024).

Potential Biases in PASSION Measurement Bias (proxy var) - Country of Origin in PASSION depending on the interpretation - should not be used for ethnicity, as this is not linked directly to the genes, see example https://medium.com/bcggamma/prac ai-responsibly-with-proxy-variable-detection-42c2156ad986

Annotator bias regarding skin tone labeling has been investigated in (Montoya et al., 2025). PASSION should evaluate its process.

used

- Measurement Bias. Measurement, or reporting, bias arises from how we choose, utilize, and measure particular features (Suresh & Guttag, 2021) (e.g. mismeasured proxy variables) (Mehrabi et al., 2021). (= e.g. someone who lives at that postal code probably has this ethnicity); –> could that be an issue with the country of origin feature?

- This study found that while using skin tone instead of race for fairness evaluations in computer vision seems objective, the annotation process remains biased by human annotators. Untested scales, unclear procedures, and a lack of awareness about annotator backgrounds and social context significantly influence skin tone labeling. This study exposes how even minor design choices in the annotation process, like scale order (dark to light instead of light to dark) or image context (face or no face, skin lesion presence), can sway agreement and introduce uncertainty in skin tone assessments. ... The researchers emphasize the need for greater transparency, standardized procedures, and careful consideration of annotator biases to mitigate these challenges and ensure fairer and more robust evaluations in computer vision. (Montoya et al., 2025) - demographic dermatology bias

- Observer bias - projecting expectations onto the research (Mester, 2022; Mester, 2017) - statistical bias

- Observer bias: When different observers view the same observation, they report it differently e.g., different observers may give differing descriptions about subtle features in the histopathology report of a skin biopsy.29 26. (Chakraborty, 2024) - dermatology

- Recall bias - respondent doesn't remember things correctly; Recall bias is another common error of interview/survey situations. It happens when the respondent doesn't remember things correctly. It's not about bad or good memory – humans have selective memory by default. After a few years (or even a few days), certain things stay and others fade. It's normal, but it makes research much more difficult. TODO: keep an eye on this when recalling evidences!! (Mester, 2022; Mester, 2017) - statistical bias

- Memory or recall bias: This is a type of bias where sufferers of a disease, often termed cases, have a greater tendency to recall a particular habit than non-sufferers, viz controls. This results in an uneven distribution of risk factors between the cases and controls. An example of this would be a case-control study to evaluate the association between dental amalgam use and the development of oral lichen planus. Those with lichen planus are more likely to recall a history of dental amalgam use than those who do not have the disease. This difference in recall between a diseased cohort and control has resulted in difficulties in assessing the association between diet and many dermatological diseases – like milk and chocolate consumption and acne, fatty meals and psoriasis, sugary meals and psoriasis, agricultural exposure to insecticides and pemphigus and so on.3–6 2. (Chakraborty, 2024) - dermatology

Research Biases

TODO: consider to move at beginning / out of data biases Researchers and their processes can also be biased in multiple ways:

- **Funding / Sponsorship bias**, when a study is deliberately supporting those findings, which the sponsor expects (**c22**; Chakraborty, 2024; Mester, 2017)

- **Data dredging bias**. The statistical methods and model are chosen to provide a certain p-value, to improve the probability of the research hypothesis being true. TODO: consider to move this to an own reporting section (Chakraborty, 2024)

- **Hypothetical bias**. Hypothetical questions lead to responses that do not reflect, what interviewees would do in real life. (**c31**; **c28**; Chakraborty, 2024) TODO: isn't this a user bias instead?

Potential Biases in PASSION   Since the PASSION dataset is already published, the research biases might already be introduced. It is not feasible during the duration of this thesis to make an evaluation on those biases. Instead, I would recommend the PASSION team and researcher in general, to check the list above carefully and take measures against them. Maybe, an external evaluation could help to detect and prevent those biases even better.

used

- Funding bias (Mester, 2022; Mester, 2017) - statistical bias

- Industry sponsorship bias: This has now been reclassified as conflict-of-interest bias. In short, the study deliberately supports the findings expected from it by its sponsors. 22.(Chakraborty, 2024) - dermatology

  Reporting biases

- Data dredging bias: It is an entirely avoidable bias. This is subdivided into two types – Fishing type and "P-value hacking" type. It involves using multiple statistical methods to get the desired p-value and selecting the statistical model that gives the p-value the author wants. This is "lamentably common" in dermatological research.16 To detect data dredging bias, always perform a "p-curve analysis" while performing a meta-analysis.17,18 Much emphasis is nowadays given to the confidence interval instead of the p-value, which gives an approximate idea of the range in which one can be 95% (or 90%, depending on the confidence interval chosen) sure that the result is correct. The confidence interval remains unaffected by p-value dredging. This subject has been reviewed in depth in recent works.18,19 15.(Chakraborty, 2024)

- Hypothetical bias: Many dermatological researches (and some life quality questionnaires like vitiQoL) use hypothetical questions – like "What would you do when some stranger asks you about your lesion?". The responses to these questions by the study participants often do not tally with what they would do in real life. This is called hypothetical bias and is avoided by adopting the ex-ante approach.31 28. (Chakraborty, 2024) - dermatology

Feature Representation Biases

Some of those biases are:

- **Omitted Variable Bias** arises when variables are not included in the model, which leads to situations for which the model is not ready for (Mehrabi et al., 2021; Mester, 2022)(Clarke, 2005; Riegg, 2008)(Mustard, 2003).

- **Collider Bias** Two variables can influence a common third variable, the collider variable. When sampling is restricted by this collider variable, it could lead to a distortion (**c4**; **c8**; **c9**; Chakraborty, 2024).

Potential Biases in PASSION The ethnicity is omitted in the PASSION dataset which could lead to issues See the medical section for more specific collider bias, maybe there could be others
used

- Omitted Variable Bias. Omitted variable bias4 occurs when one or more important variables are left out of the model (Clarke, 2005; Riegg, 2008)(Mustard, 2003). Something that the model was not ready for(Mehrabi et al., 2021). did not take into account (Mehrabi et al., 2021)

- Omitted variable bias (Mester, 2022; Mester, 2017) - statistical bias

- Collider Bias: This is an under-appreciated bias, and often confused with a confounder. This is especially seen in observational studies where it is defined as a distortion produced by the restriction of sampling by a collider variable. A collider variable is defined as one that has an independent effect on the outcome studied apart from the studied variable. In simpler terms, collider bias occurs when exposure and development influence a common third variable. That variable or collider is controlled by study design or in the analysis. An example is the observation that psoriasis patients tend to have more depression and anxiety disorders. Since severe psoriasis patients tend to get hospitalised and also get screened for mental health issues, a spurious association between them could have been obtained due to collider bias. The two variables viz psoriasis and depression converged, i.e., collided, into a single outcome – hospitalization.8,9 4. (Chakraborty, 2024) - dermatology

Imaging Biases

Dealing with images can lead to a whole other set of challenges, which can lead to biases. The challenges are for example technical variations in hardware and software but also differences in how images are gathered or what is in it (Young et al., 2020).

Those biases can be introduced :

- **Image Quality Bias**. The quality of an image (zoom level, focus, lightning) could be associated with the classification (Young et al., 2020)

- **Visual Artifact Bias**. Other artifacts, such as presence of hair or surgical ink markings on dermatology images, can decrease classification performance (**Winkler et al.**; **2019 & Bisla et al.**; **2019 (from Young_2020)**)

- **Field of View Bias**. What view is captured in the image can interfere with prediction quality what is it, consequence (**Mishra et al.**; **2019 from Young_2020**)

Potential Biases in PASSION   The PASSION model could learn to associate unrelated visual effects, hair, body parts or image quality with a disease.
used

- Image quality. Several barriers to AI implementation in the clinic need to be overcome with regards to imaging (Figure 1). These include technical variations (e.g., camera hardware and software) and differences in image acquisition and quality (e.g., zoom level, focus, lighting, and presence of hair). For example, the presence of surgical ink markings is associated with decreased specificity (Winkler et al., 2019), field of view can significantly affect prediction quality (Mishra et al., 2019), and classification performance improves when hair and rulers are removed (Bisla et al., 2019). We have developed a method to measure how model predictions might be biased by the presence of a visual artifact (e.g., ink) and proposed methods to reduce such biases (Pfau et al., 2019). Poor quality images are often excluded from studies, but the problem of what makes an image adequate is not well studied. Ideally, models need to be able to express a level of confidence in a prediction as a function of image quality and appropriately direct a user to retake photos if needed. (Young et al., 2020) - dermatology

Medical Biases

In ML for health care, there are special medical versions of the mentioned biases as well as completely new biases. They require special attention, since they directly influence the diagnosis or treatment of a disease.

Those biases can be introduced:

- **Berkesonian bias** occurs in hospital-based studies when two variables influence hospital or clinical attendance independently. This can lead to a distorted estimation of the relationship between those variables because the study population of hospitalized patients is not representative of the whole population (**c3**; **c7**; Chakraborty, 2024)

- **Informed presence bias**, the probability to get screened for other diseases is higher for people who seek medical care. Like Berkesonian bias, this can lead to misleading interpretations of relationships between two diseases (**c27**; **c23**; Chakraborty, 2024)

- **Diagnostic access bias**, depending on the geographical location, individuals have better access to medical care. Therefore, their disease prevalence could appear to be higher and diseases could be diagnosed earlier. (**c19-c21**; Chakraborty, 2024)

- **Diagnostic reference test bias** is a **verification bias**, where not all individuals receive the same reference test for the diagnostic process, potentially leading to different diagnoses. (**c21**; Chakraborty, 2024)

- **Mimicry bias**, exposures to treatment options can cause a disease which presents itself similar to the study disease, which potentially creates misleading data (**c28**; **c25**; Chakraborty, 2024)

- **Unacceptable Disease bias**. When a disease is socially unacceptable, it can result in under-reporting of the same disease (**c30**; **c27**; Chakraborty, 2024)

- **Healthy volunteer selection bias**, is a type of self-selection bias where the volunteers are in general healthier than the population due to more interest in health (Delgado-Rodríguez & Llorca, 2004)

Potential Biases in PASSION   Berkesonian bias depending on the chosen hospitals Informed presence bias regarding correlation between impedigo and the other diseases Diagnostic access bias can somewhat be addressed by PASSION, since its dataset includes samples of later states of diseases. However, in the PASSION context itself, this bias could still be relevant. Diagnostic reference test bias could be inherited in the PASSION dataset, depending on how the dermatologists work. Mimicry bias is not relevant regarding the exposures since PASSION does not hold any exposure data. However, diseases which mimicry others could lead to issues if they are not detected.

used

- Berkesonian Bias: Named after Dr. Joseph Berkeson, this bias reflects the variation in rates of hospital admission or clinic attendance for different diseases. For example, if a study is conducted to know the effect of pregnancy on syphilis in an antenatal clinic, we are likely to get biased data since the two conditions, viz pregnancy and syphilis, are both likely to affect clinic attendance and all observations related to the relationship between pregnancy and syphilis.7 3. (Chakraborty, 2024) - dermatology

- Informed presence bias: Simply, a person attending a health center is more likely to get screened for other unrelated comorbidities than those not attending a health center e.g., the finding psoriasis is associated with depression has now been criticised because those having psoriasis also have a greater chance to be screened for depression since they are already attending a health center.27 23. (Chakraborty, 2024) - dermatology

- Diagnostic Access Bias: Individuals in certain geographical localities have better access to medical care and, hence, may appear to have higher disease prevalence. For example, atopic dermatitis is believed to be commoner in the West – this could be due to better and earlier diagnostic facilities available than in India.19,20 17.(Chakraborty, 2024)

- Diagnostic reference test bias: These bias results when all individuals do not receive the same reference test. e.g., direct immunofluorescence studies may not be done for all patients with pemphigus vulgaris some patients

may receive only a skin biopsy-based diagnosis. It is a subtype of verification bias. Another variation of this type of bias is partial reference bias, where only some of the study participants receive the index and the reference tests.21(Chakraborty, 2024)

- Mimicry bias: When an exposure causes a disease that resembles the study disease, mimicry bias can result. For example, certain drugs are known to cause a pityriasis rosea-like reaction, which, although looks like pityriasis rosea, differs from it.28 25.(Chakraborty, 2024) - dermatology

- Unacceptable disease bias: This occurs in socially unacceptable diseases like leprosy and STDs, which result in under-reporting.30 27. (Chakraborty, 2024) - dermatology

- TODO: Other such studies were conducted in [(Fry et al., 2017)] which states that UK Biobank, a large and widely used genetic dataset, may not represent the sampling population. Researchers found evidence of a "healthy volunteer" selection bias. [150] has other examples of studies on existing biases in the data used in the medical domain. [157] also looks at machine-learning algorithms and data utilized in medical fields, and writes about how artificial intelligence in health care has not impacted all patients equally.(Mehrabi et al., 2021) –> [150] also provides an ovverview over the impact of social determinants on health, such as Economic stability, neighborhood and physical environment, education, food, community and scial context, access to healthcare and quality

- The healthy volunteer effect is a particular case: when the participants are healthier than the general population. (Delgado-Rodríguez & Llorca, 2004)

Temporal Biases

Differences in populations and their behaviour over time can lead to temporal biases (Olteanu et al., 2019). Certain studies require to track temporal data, to learn about their behaviour over time. Disease progression is also a factor measured over time (Mehrabi et al., 2021). For PASSION, temporal biases are currently irrelevant, since PASSION contains images independently of time and is not tracking the disease progression. Therefore, the listed biases in this chapter are not explained in detail, refer to the sources for further information.

Examples for temporal data biases are:

- **Longitudinal Data Fallacy** (Mehrabi et al., 2021)

- **Chronological bias** (**c9**; **c13**; Chakraborty, 2024)

- **Immortal time bias** (**c24**; **c20**; Chakraborty, 2024)

2.2.9 Algorithmic Biases

When an algorithm adds biases to unbiased input data one speaks of **Algorithmic Bias** (Baeza-Yates, 2018). This could occur due to algorithmic design choices like optimization functions, regularizations and statistically biased estimators (Danks & London, 2017).

used

- Algorithmic Bias. Algorithmic bias is when the bias is not present in the input data and is added purely by the algorithm (Baeza-Yates, 2018). The algorithmic design choices, such as use of certain optimization functions, regularizations, choices in applying regression models on the data as a whole or considering subgroups, and the general use of statistically biased estimators in algorithms (Danks & London, 2017), can all contribute to biased algorithmic decisions that can bias the outcome of the algorithms.(Mehrabi et al., 2021).

User Algorithm Interaction Biases

- **User Interaction Bias**. This biases can be triggered by the user interface or the user themselves. The user interface influences the user to behave in a certain way, which could introduce bias in the user behaviour. Users impose this (or their own) biased behavior through interaction on the algorithm (Baeza-Yates, 2018). **Presentation bias** and **Ranking bias** are further subtypes mentioned by Lerman and Hogg (2014) and Mehrabi et al. (2021).

- **Emergent Bias**. When real users interact with an algorithm, this bias arises some time after the design was completed due to changes in population. It appears more likely in user interfaces (Friedman & Nissenbaum, 1996).

Potential Biases in PASSION   The user interaction biases, especially the emergent bias could potentially become an issue for PASSION, when the project starts to become publicly available teledermatology. Also, the interface design should be evaluated, so that no presentation or ranking bias gets introduced.

used

- Emergent Bias. Emergent bias occurs as a result of use and interaction with real users. This bias arises as a result of change in population, cultural values, or societal knowledge usually some time after the completion of design (Friedman & Nissenbaum, 1996). This type of bias is more likely to be observed in user interfaces, ... This type of bias can itself be divided into more subtypes, as discussed in detail in (Friedman & Nissenbaum, 1996). (Mehrabi et al., 2021). probably less relevant at the first stage

- User Interaction Bias. User Interaction bias is a type of bias that can not only be observant on the Web but also get triggered from two sources—the user interface and through the user itself by imposing his/her self-selected

biased behavior and interaction (Baeza-Yates, 2018). This type of bias can be influenced by other types and subtypes, such as presentation and ranking biases. (Mehrabi et al., 2021). – more relevant for later, when the application would become bigger

- Presentation Bias. Presentation bias is a result of how information is presented (Baeza-Yates, 2018) (can only click on content they see, could be the case that user does not see all info on web) (Mehrabi et al., 2021).

- Ranking Bias. The idea that top-ranked results are the most relevant and important will result in attraction of more clicks than others. This bias affects search engines (Baeza-Yates, 2018) and crowdsourcing applications (Lerman & Hogg, 2014).(Mehrabi et al., 2021).

External Influence Biases

Those biases can be introduced :

- **Evaluation Bias**. When inappropriate or disproprtionate benchmarks are used in model evaluation, they can introduce the benchmarks biases into the model. (Suresh & Guttag, 2021; Buolamwini & Gebru, 2018)

- **Incorporation bias**. When index tests in diagnostic accuracy studies are part of the reference tests, this results in elevated sensitivity for the index tests (**c21**; **c25**; **c26**; Chakraborty, 2024; Young et al., 2020).

- **Popularity Bias**. More popular items tend to be exposed more. Popularity metrics can be manipulated though or not reflecting good quality, this can lead to bias (Ciampaglia et al., 2018; Mehrabi et al., 2021).

- **Generalization Issues**. (**<empty citation>**) TODO: add those from young

Potential Biases in PASSION   TODO: add used

- Evaluation Bias. Evaluation bias happens during model evaluation (Suresh & Guttag, 2021). This includes the use of inappropriate and disproportionate benchmarks for evaluation of applications such as Adience and IJB-A benchmarks. These benchmarks are used in the evaluation of facial recognition systems that were biased toward skin color and gender (Buolamwini & Gebru, 2018), and can serve as examples for this type of bias (Suresh & Guttag, 2021). (Mehrabi et al., 2021). – important for this thesis

- Incorporation bias: This is principally relevant for diagnostic accuracy studies when the index test forms a part of the reference test, resulting in elevated sensitivity e.g., if one wants to compare the grattage test vs. dermoscopy

in psoriasis and does dermoscopy only from areas of grattage positivity, one would get a very high sensitivity for the grattage test because it was incorporated into the reference test, i.e., dermoscopy.25,26 21.(Chakraborty, 2024)

- Popularity Bias. Items that are more popular tend to be exposed more. However, popularity metrics are subject to manipulation—for example, by fake reviews or social bots (Ciampaglia et al., 2018). ... this presentation may not be a result of good quality; instead, it may be due to other biased factors. (Mehrabi et al., 2021).

2.2.10 User Biases

Cognitive Biases

Biases which are related to human perception belong to the category of cognitive biases. They are affecting how data should be presented and interpreted (Mester, 2017)

Those biases can be introduced :

- **Confirmation Bias**. When people have pre-conceptions, they will only listen to the part of presented information which reinforce those "facts", regardless whether the facts are true or not (Mester, 2017). In health-care, this can be observed when patients report increases in diseases due to potentially nonfactual information they found on the internet (**c15**; **c14**; Chakraborty, 2024).

- **Belief Bias**. A stronger version of the confirmation bias: Someone who is affected by this bias is so sure about their own gut feelings that they will ignore results of a data research project (Mester, 2017).

- **Previous Opinion Bias**. When performing multiple tests, the knowledge about the outcome of the previous tests probably influences the results (Chakraborty, 2024)

- **Cause-Effect Bias**. The famous senctence "correlation does not imply causation" can be used here - when correlation between two variables is misinterpreted as a cause-effect in the wrong direction, this bias applies (Mester, 2017)

- **Historical Bias**. Preexisting biases in the world can affect the data generation process (Suresh & Guttag, 2021). Even if they reflect the current reality, it is worth to consider whether those biases should affect the algorithms in question (Mehrabi et al., 2021).

- **Content Production Bias**. User generated contents can introduce biases by systematical differences in the production process, stucture and appearance, which might stem from the users background (Olteanu et al., 2019).

Potential Biases in PASSION   For PASSION, confirmation bias could lead to issues in the initial diagnosis and could therefore lead to biased data labeling. Same with the previous opinion bias. The later can be reduced when it is ensured, that the labeling experts are diagnosing the diseases independently of each other, so that they do not know the previous opinions. Cause-Effect bias is lesser an issue for PASSION, since the causes of the diseases are not analyzed. It could more be an inherit problem, that the algorithm learns wrong causes for diseases, such as appearing hair Historical bias can affect PASSIONs process in various ways. In PASSION context, Content Production Bias could have an impact on how the images are taken.

used

•

- Cognitive bias (Mester, 2017) - statistical bias

- Previous opinion bias: In performing a second diagnostic test, if the result of a previous test is known, it is likely to influence the result. An extension of this is the Greenwald's law of lupus: the Sontheimer amendment – anything and everything that happens to a lupus erythematosus patient is correctly or incorrectly attributed to lupus.32 29. (Chakraborty, 2024) - dermatology

- Confirmation bias: This bias occurs when study participants have a preconceived notion of their disease that may not be based on facts. For example, we have observed that in North India many tinea patients report an increase in their disease due to taking meat, fish, and other so-called "hot foods". They may also present information they have collected from the internet which reinforces their beliefs.15 14.(Chakraborty, 2024) - dermatology

- Cause-effect bias (Mester, 2022; Mester, 2017) - statistical bias

- Historical Bias. Historical bias is the already existing bias and socio-technical issues in the world and can seep into from the data generation process even given a perfect sampling and feature selection (Suresh & Guttag, 2021). ... search results were of course reflecting the reality, but whether or not the search algorithms should reflect this reality is an issue worth considering (Mehrabi et al., 2021) - maybe relevant

- Content Production Bias. Content Production bias arises from structural, lexical, semantic, and syntactic differences in the contents generated by users (Olteanu et al., 2019). (Mehrabi et al., 2021) – could the quality of the pictures been related to this as well?

Behavioral Biases

Those biases can be introduced :

- **Behavioral Bias**. User behaviour can differ depending on the platforms, contexts, cultures, or datasets (Olteanu et al., 2019).

- **Self-Selection Bias**. This subtype of selection bias occurs when study participants can select themselves. Less proactive people, people with less time or interest will be excluded or underrepresented (Mester, 2022; Mehrabi et al., 2021). **Non-Responder bias** is a subtype, where part of the population is not responding e.g. to fill out a survey or post-study responses queried by postal services (Chakraborty, 2024). TODO: maybe categorize this in the data biases or the healthy volunteer bias here

- **Social Bias**. When the actions of others affect our judgment, it is called social bias. For example ratings in juries can be affected by this (Baeza-Yates, 2018).

Potential Biases in PASSION   For PASSION the behavioral biases can affect who is going to the dermatologists for what reasons. Therefore, the approach to use data from different countries may be benefitial, since potentially the cultural differences could differ. Self-selection is an issue, since only those patients can be included in the database which go to the hospitals.

used

- Self-Selection Bias. Self-selection bias4 is a subtype of the selection or sampling bias in which subjects of the research select themselves. (Mehrabi et al., 2021)

- Self-selection bias - when you let the subjects of the analyses select themselves, less proactive people will be excluded TODO: could be an issue as well for PASSION, couldn't it? since the doctors probably ask the clients. One way to go is to default should be to provide access to the data. but is it ethical? (Mester, 2022; Mester, 2017)- statistical bias A variation of this is non-responder bias, where non-responders to a questionnaire differ significantly from responders.9 9. (Chakraborty, 2024) - dermatology

- Social Bias. Social bias happens when others' actions affect our judgment (Baeza-Yates, 2018). (case where we want to rate or review an item with a low score, but when influenced by other high ratings, we change our scoring thinking that perhaps we are being too harsh [(Baeza-Yates, 2018), (Wang & Wang, 2014).) (Mehrabi et al., 2021)

- Behavioral Bias. Behavioral bias arises from different user behavior across platforms, contexts, or different datasets (Olteanu et al., 2019). (Mehrabi et al., 2021) maybe, people from different countries go to the dermatologist for different diseases, based on cultural differences?

Publication Biases

Those biases can be introduced :

- **Publication Bias**

- **Hot stuff bias** is a subtype of publication bias, where Journals are less critical about trending topics, which lead to more frequent publishing of those topics. This in turn can lead to flawed meta-analyses regarding those topics (**c22**; **c23**; **c19**; Chakraborty, 2024).

- **All is Well Bias**. This bias is a different view on the hot stuff bias. Theories which align with the view of the majority are more likely to be pubhlished than an opposing view (**c7**; **c10-12**; Chakraborty, 2024).

- **Rethoric Bias**. Charismatic writing or when the press is more vocal about findings can lead to greater influence over individuals than other available facts (Chakraborty, 2024).

- **Novelty Bias**. Newer interventions appear to be better. Over time, this effect decreases (Chakraborty, 2024).

Potential Biases in PASSION   These biases are relevant for all researchers. They should kept in mind when interpreting, publishing and peer-reviewing papers. used

- Hot stuff bias: Editors of journals may be less critical about topics that are "fashionable" or currently in vogue and consequently end up publishing them more frequently, resulting in publication bias as well as hot stuff bias. It can result in flawed meta-analyses based on these studies. An example is how cutaneous manifestations of COVID-19 were published. Indian Journal of Dermatology Venereology and Leprosy stood out by choosing not to publish anything and everything related to COVID-19, thus reducing hot stuff bias.22,23 19. (Chakraborty, 2024)

- All is well bias: It is a subjective bias where theories supported by the majority tend to get more easily published than the opposing view supported by the minority. For example, ideas on the origin of endemic pemphigus supporting autoimmunity are more likely to be published than theories exploring an infectious trigger. According to some authors, this bias is very difficult to eliminate and is a variant of publication bias.10-12 7.(Chakraborty, 2024) - dermatology

- Rhetoric bias: A more charismatic piece of writing has a greater influence on the study participants than other available literature. An example is the wider use of sunscreen for polymorphous light eruption over photoprotective strategies like umbrellas, broadbrimmed hats, etc, because the lay press is more vocal about sunscreens.14 11. (Chakraborty, 2024) - dermatology

- Novelty bias: The newer an intervention, the better it appears, and with time, its efficacy seems to decrease. When ligelizumab, an IgE antagonist was first discovered, ligelizumab was believed to be better than omalizumab; however, evidence soon pointed to the contrary. 16.(Chakraborty, 2024) - dermatology

Medical Biases

Those biases can be introduced :

- **Popularity Bias**. In medicine, when more popular diseases (usually well-known or stigmatized ones) get compared with less popular diseases, clinic rates can show a distorted view. The more popular diseases appear to be over-represented over more commoner ones (**c9**; **c6**; Chakraborty, 2024).

- **Apprehension Bias**. Fear related to an upcoming procedure can lead to false evaluations, e.g. when measuring blood pressure (**c13**; Chakraborty, 2024).

- **Hawthrone bias**. Subjects might modify their behaviour when they know they are being watched. This bias can be practically utilized by introducing regular follow-ups (**c8**; Chakraborty, 2024).

- **Centripetal Bias**. Better reputations affect to which physicians or hospitals patients tend to go to. Famous specialists probably see more cases in regards of their specialty than others (**c12**; Chakraborty, 2024).

Potential Biases in PASSION PASSION must be careful in interpreting the metadata. Since the data is from hospitals, they could be biased towards more popular diseases. PASSION can potentially use Hawthrone bias to improve the work of the annotators. Centripetal bias can also be used when selecting the partners to work with.

used

- Popularity Bias: This bias arises when a particular disease is more popular (i.e. either more well-known or more stigmatised) among the participants than the disease with which it is compared. For example, if a study compares clinic attendance rates among various dermatological disorders, one would see vitiligo patients are over-represented over melasma. While melasma is commoner in the normal population, vitiligo, due to its popularity because of media publicity and other factors, tends to present earlier.9 6. (Chakraborty, 2024) - dermatology

- Apprehension bias: This results from fear and apprehensions related to an impending procedure. The classic example is the false elevation of blood pressure because the person is apprehensive of his or her blood pressure being measured.13 A variant of this is the Hawthorne bias, where subjects modify their behavior, such as regularly taking a prescribed drug or exercising, simply because they know they are being watched, but not due to any apprehensions. Hawthorne bias is practically utilised in many leprosy clinics since regular follow-up has been shown to improve adherence to therapy based on Hawthorne bias. 8. (Chakraborty, 2024) - dermatology

- Centripetal bias: Patients tend to go to more reputed physicians and hospitals than others. For example, a famous or better-known cosmetologist

with a good reputation tends to see more cases than other cosmetologists. 12.(Chakraborty, 2024) - dermatology

2.2.11 Sensitive Features

Research showed sensitive features which are prone to bias. Those led to issues in existing AI applications and their usage should therefore be investigated carefully.

Table 2.4 summarizes mentioned features. Since PASSION tries to improve the classification of skin diseases with photographs alone, those features which directly influence the appearance of a disease on the skin are most important. Skin type including undertone are affecting the diseases appearance, but also genetic factors, gender and the age TODO: doublecheck impact of age of a patient can influence how the diseases manifest themselves. Therefore, they are directly relevant for skin detection and must be considered in the dataset.

Other demographic factors can have an impact on the progression of skin diseases or increase the prevalence of skin conditions (e.g. tropical vs dry climates) TODO: add source. While this information could support the diagnostic process, it is not visible in the images. Therefore, they could potentially be relevant and enhance the diagnostic process when included into the PASSION dataset. However, the linkage is weak and could introduce further biases.

For completeness, Table 2.4 shows further sensitive demographic features which seem not to be linked to dermatology according to research.

TODO: check what to do with those additional features: Other important features according to ((Montoya et al., 2025) 13): lesion type, anatomical location of lesion, img characteristics such as source, imaging techniques, resolution, real vs. artificially generated

Bias-Sensitive Features	Mentioned in Context of	
	ML	**Dermatology**
Dermatology Related Features		
Skin Type	X[1,2,7]	X[12,13]
Skin Undertones		X[13]
Demographic Features		
Relevant for Skin Disease Detection		
Age	X[7,11]	X[13]
Gender/Sex	X[1,2,7,8,9,10,11]	X[13]
Gender and Skin Type Subgroups	X[1,2]	
Ethnicity/Race	X[1,2,4,5,6,7,11]	X[12,13]
Potentially Relevant for Skin Disease Detection		
Geographic Location	X[1,3]	
Socio-Economic Status	X[6]	X[12]
Disabilities	X[7,11]	
Not Relevant for Skin Disease Detection		
Familial status	X[7]	
Marital status	X[7,11]	
Nationality/National origin	X[7,11]	
Recipient of public assistance	X[7]	
Religion	X[7,11]	

[1] (Mehrabi et al., 2021) [6] (Vickers & Fouad, 2014) [11] (Hajian & Domingo-Ferrer, 2013)
[2] (Buolamwini & Gebru, 2018) [7] (Chen et al., 2019)
[3] (Shankar et al., 2017) [8] (Zhao et al., 2017) [12] (Young et al., 2020)
[4] (Manrai et al., 2016) [9] (Bolukbasi et al., 2016) [13] (Montoya et al., 2025)
[5] (Fry et al., 2017) [10] (Zhao et al., 2018)

Table 2.4: Features which often hold biases

TODO: decide which table to use, more or less extensive citations?

Bias-Sensitive Features	Mentioned in Context of	
	ML	Dermatology
Dermatology Related Features		
Skin Type	X[1,3]	X[5,6]
Skin Undertones		X[6]
Demographic Features		
Relevant for Skin Disease Detection		
Age	X[3,4]	X[6]
Gender/Sex	X[1,3,4]	X[6]
Gender and Skin Type Subgroups	X[1]	
Ethnicity/Race	X[1,2,3,4]	X[5,6]
Potentially Relevant for Skin Disease Detection		
Geographic Location	X[1]	
Socio-Economic Status	X[2]	X[5]
Disabilities	X[3,4]	
Not Relevant for Skin Disease Detection		
Familial status	X[3]	
Marital status	X[3,4]	
Nationality/National origin	X[3,4]	
Recipient of public assistance	X[3]	
Religion	X[3,4]	

[1] (Mehrabi et al., 2021)   [3] (Chen et al., 2019)   [5] (Young et al., 2020)
[2] (Vickers & Fouad, 2014)   [4] (Hajian & Domingo-Ferrer, 2013)   [6] (Montoya et al., 2025)

Table 2.5: Features which often hold biases

2.3 Fairness

As Mehrabi et al. (2021) states, "in order to be able to fight against discrimination and achieve fairness, one should first define fairness". However, research is no clear agreement for a single definition of fairness. Further, the fairness definitions are also prone to biases since different cultures and preferences influence the perception of fairness. In spite of those challenges, fairness can broadly be defined as absence of bias in decision making. While fairness is very desirable, it "can be surprisingly difficult to achieve in practice" (Mehrabi et al., 2021). This chapter summarizes the existing fairness methods, which aim to achieve fairness. The fairness methods were listed by Mehrabi et al. (2021).

2.3.1 Fairness Definitions

Fairness Definitions	Mentioned in Context of	
	ML	Dermatology
Group Fairness		
Conditional Statistical Parity	X[1,3,10]	
Demographic/Statistical Parity	X[1,3,4,5]	
Equalized Odds	X[1,2,3]	
Equal Opportunity	X[1,2,3]	
Treatment Equality	X[1,7]	
Test Fairness	X[1,3,8]	
Subgroup Fairness		
Subgroup Fairness	X[1,11,12]	
Individual Fairness		
Counterfactual Fairness	X[1,5]	
Fairness Through Awareness	X[1,4,5]	
Fairness Through Unawareness	X[1,5,6]	
Not Categorized		
Fairness in Relational Domains	X[1,9]	

[1] (Mehrabi et al., 2021)   [5] (Kusner et al., 2017)   [9] (Farnadi et al., 2018)
[2] (Hardt et al., 2016)   [6] (Grgic-Hlača et al., 2016)   [10] (Corbett-Davies et al., 2017)
[3] (Verma & Rubin, 2018)   [7] (Berk et al., 2017)   [11] (Kearns et al., 2018)
[4] (Dwork et al., 2012)   [8] (Chouldechova, 2017)   [12] (Kearns et al., 2019)

Table 2.6: Fairness definitions as stated in Mehrabi et al. (2021)

TODO: deside which version should be used

Fairness Definitions	Mentioned in Context of	
	ML	**Dermatology**
Group Fairness		
Conditional Statistical Parity	X	
Demographic/Statistical Parity	X	
Equal Opportunity	X	
Treatment Equality	X	
Test Fairness	X	
Equalized Odds	X	
Subgroup Fairness		
Subgroup Fairness	X	
Individual Fairness		
Counterfactual Fairness	X	
Fairness Through Awareness	X	
Fairness Through Unawareness	X	
Not Categorized		
Fairness in Relational Domains	X	

Table 2.7: Fairness definitions as summarized in Mehrabi et al. (2021)

TODO: potential bias in (Kearns et al., 2019) in this bc same author of the algorithm tested it

As shown in the Table 2.7, there are mainly three categories for fairness definitions. According to Mehrabi et al. (2021), fairness can be achieved on a group level, subgroup level or even for a individual. Group fairness is about treating different groups as equal. Individual fairness tries to achieve similar predictions for similar individuals. Subgroup fairness tries to incorporate the best properties of the other two levels to improve the outcome in larger collections of subgroups (Mehrabi et al., 2021).

The specific fairness definitions can be found in Mehrabi et al. (2021). In general, they try to get similar probability outcomes for 'unprotected' or 'protected' groups. This list summarizes how they work:

- **Demographic/Statistical Parity** and **Conditional Statistical Parity**: The parity checks that likelihood of a positive outcome is the same for both protected groups (Mehrabi et al., 2021). The conditional version adds legitimate factors before calculating the statistical parity (Corbett-Davies et al., 2017).

- **Equalized Odds**, **Test Fairness**. and **Equal Opportunity**: All those methods protected and unprotected groups should have equal rates for a positive outcome when belonging to a positive class, essentially comparing the true positive rates. **Equalized Odds** is a more restrictive since it also checks for similar false positive rates (Verma & Rubin, 2018; Mehrabi et al., 2021). TODO: add M48 and M87 for demographic parity

- **Treatment Equality**: It compares the false negative and false positive rates (Wang & Wang, 2014)

- **Counterfactual Fairness**: This approach is different from the others as it is testing the same individual in both different demographic groups with the intention that the outcome is the same (Kusner et al., 2017; Mehrabi et al., 2021). It differs from to the first group of fairness metrics since it does not compare the likelihoods of the outcomes for any person in a group, but checking how the exact same individual would be treated if it was in the other group.

- **Fairness Through Awareness**: This method compares similar individuals based on similarity metrics to get a similar outcome (Mehrabi et al., 2021)

- **Fairness Through Unawareness**: This measure is ensuring that protected attributes are not explicitly used in desicion-making (Grgic-Hlača et al., 2016; Kusner et al., 2017).

- **Fairness in Relational Domains**: This notion also takes into consideration relational structures between individuals (Farnadi et al., 2018).

2.3.2 Challenges of Fairness Definitions

2.4 Mitigation Methods

see text from bias chapter - Further, AI engineers need to know what prevention methods are available to reduce the biases (Mehrabi et al., 2021).

2.5 Extensive Sources

2.5.1 Discrimination vs. Biases

- bias and discrimination = source of unfairness. Discrimination can be considered as a source for unfairness that is due to human prejudice and stereotyping based on the sensitive attributes, which may happen intentionally or unintentionally, while bias can be considered as a source for unfairness that is due to the data collection, sampling, and measurement. Although bias can also be seen as a source of unfairness that is due to human prejudice and stereotyping, in the algorithmic fairness literature it is more intuitive to categorize them as such according to the existing research in these areas. In this survey, we mainly focus on concepts that are relevant to algorithmic fairness issues. (Mehrabi et al., 2021)

- Explainable Discrimination. Differences in treatment and outcomes amongst different groups can be justified and explained via some attributes in some cases. In situations where these differences are justified and explained, it is not considered to be illegal discrimination and hence called explainable

[77]. In [77], authors present a methodology to quantify the explainable and illegal discrimination in data. They argue that methods that do not take the explainable part of the discrimination into account may result in non-desirable outcomes, so they introduce a reverse discrimination which is equally harmful and undesirable. They explain how to quantify and measure discrimination in data or a classifier's decisions which directly considers illegal and explainable discrimination.(Mehrabi et al., 2021)

- Unexplainable Discrimination. In contrast to explainable discrimination, there is unexplainable discrimination in which the discrimination toward a group is unjustified and therefore considered illegal. Authors in [77] also present local techniques for removing only the illegal or unexplainable discrimination, allowing only for explainable differences in decisions. These are preprocessing techniques that change the training data such that it contains no unexplainable discrimination. We expect classifiers trained on this preprocessed data to not capture illegal or unexplainable discrimination. Unexplainable discrimination consists of direct and indirect discrimination.(Mehrabi et al., 2021)

    - Direct Discrimination. Direct discrimination happens when protected attributes of individuals explicitly result in non-favorable outcomes toward them [164]. ... these traits that are considered to be "protected" or "sensitive" attributes in computer science literature (Mehrabi et al., 2021)
    - Indirect Discrimination. In indirect discrimination, individuals appear to be treated based on seemingly neutral and non-protected attributes; however, protected groups, or individuals still get to be treated unjustly as a result of implicit effects from their protected attributes (Mehrabi et al., 2021)

- Discrimination can be either direct or indirect. Direct discrimination occurs when decisions are made based on sensitive attributes. Indirect discrimination occurs when decisions are made based on nonsensitive attributes which are strongly correlated with biased sensitive ones (Hajian & Domingo-Ferrer, 2013)

- ML algorithms reflect cognitive biases in humans, which can result in predictions and decisions that are "unfair" [5]. Unfairness in ML decision-making equates to prejudice or favoritism based on inherent or acquired characteristics of an individual or group [6]. Within medical research, the term bias refers to "a feature of the design of a study, or the execution of a study, or the analysis of the data from a study, that makes evidence misleading" [7,8]. (**Montonaya_2025**)

2.5.2 Bias Introduction

- compared SAVRY, a tool used in risk assessment frameworks that includes human intervention in its process, with automatic machine learning methods

in order to see which one is more accurate and more fair. Conducting these types of studies should be done more frequently, but prior to releasing the tools in order to avoid doing harm (Mehrabi et al., 2021).

- **Assessment Tools** An interesting direction that researchers have taken is introducing tools that can assess the amount of fairness in a tool or system. For example, Aequitas [136] is a toolkit that lets users to test models with regards to several bias and fairness metrics for different population subgroups. Aequitas produces reports from the obtained data that helps data scientists, machine learning researchers, and policymakers to make conscious decisions and avoid harm and damage toward certain populations. AI Fairness 360 (AIF360) is another toolkit developed by IBM in order to help moving fairness research algorithms into an industrial setting and to create a benchmark for fairness algorithms to get evaluated and an environment for fairness researchers to share their ideas [11]. These types of toolkits can be helpful for learners, researchers, and people working in the industry to move towards developing fair machine learning application away from discriminatory behavior (Mehrabi et al., 2021).

- At first sight, automating decisions may give a sense of fairness: classification rules do not guide themselves by personal preferences. However, at a closer look, one realizes that classification rules are actually learned by the system (e.g., loan granting) from the training data. If the training data are inherently biased for or against a particular community (e.g., foreigners), the learned model may show a discriminatory prejudiced behavior. In other words, the system may infer that just being foreign is a legitimate reason for loan denial. (Hajian & Domingo-Ferrer, 2013)

- One might think of a straightforward preprocessing approach consisting of just removing the discriminatory attributes from the data set. Although this would solve the direct discrimination problem, it would cause much information loss and in general it would not solve indirect discrimination. (Hajian & Domingo-Ferrer, 2013)

- Hence, there are two important challenges regarding discrimination prevention: one challenge is to consider both direct and indirect discrimination instead of only direct discrimination; the other challenge is to find a good trade-off between discrimination removal and the quality of the resulting training data sets and data mining models. (Hajian & Domingo-Ferrer, 2013)

- Most AI systems and algorithms are data driven and require data upon which to be trained. Thus, data is tightly coupled to the functionality of these algorithms and systems. In the cases where the underlying training data contains biases, the algorithms trained on them will learn these biases and reflect them into their predictions. As a result, existing biases in data can affect the algorithms using the data, producing biased outcomes. Algorithms can even amplify and perpetuate existing biases in the data.(Mehrabi et al., 2021).

- In addition, algorithms themselves can display biased behavior due to certain design choices, even if the data itself is not biased. The outcomes of these biased algorithms can then be fed into real-world systems and affect users' decisions, which will result in more biased data for training future algorithms.(Mehrabi et al., 2021).

- Bias can exist in many shapes and forms, some of which can lead to unfairness in different downstream learning tasks. In (Suresh & Guttag, 2021), authors talk about sources of bias in machine learning with their categorizations and descriptions in order to motivate future solutions to each of the sources of bias introduced in the paper. In (Olteanu et al., 2019), the authors prepare a complete list of different types of biases with their corresponding definitions that exist in different cycles from data origins to its collection and its processing.(Mehrabi et al., 2021).

- The list of biases that can occur in any research is considerably long, and certainly not all of them can be avoided. However, dermatologists should be well aware of them because: 1. While conducting systematic reviews and metaanalysis, PRISMA guidelines need to be followed, and the PRISMA checklist requires a very exhaustive list of declarations to be made by the authors as to how biases in the individual studies were detected, their types and whether they were included in the systematic review or not. This usually requires more than two authors working independently. 2. Biases can result in dramatically opposite inferences, which may not be biologically plausible; the knowledge of the biases can help detect them and thereby negate such findings. 3. The knowledge of biases is a vital part of the postgraduate dermatology curriculum and is a mustknow area for thesis/dissertation purposes. 4. Since there is no fool-proof way to avoid all biases, the help of a well-qualified biostatistician might help detect and prevent many of these biases in research. (Chakraborty, 2024)

Already rewritten: The following categorization was modeled with the intent to show that the different biases are intertwined and one should consider the effects between each other in the cycle to address them correctly (Mehrabi et al., 2021)

2.5.3 Biases Extensive Sources

Data Biases

Data biases (data to algorithm (biases in data which might have an impact on biased algorithmic outcomes (Mehrabi et al., 2021)))

-

Algorithmic Biases

Algorithmic biases (Algorithm to user (A modulates U behaviour, biases in algorithm might lead to introduce biases in user behaviour and affect it as a consequence)) (Mehrabi et al., 2021)

-

User Biases

User to Data (user-generated data, inherent biases in users could be reflected in the data they generate; biases in last section might introduce further bias in this process) (Mehrabi et al., 2021)

-

Dermatology Biases

- Equity. AI has the potential to worsen health-care disparities, as recognized by the popular media (Khullar, 2019), particularly in dermatology (Adamson and Smith, 2018). The first concern is adequate representation of underserved populations in training data. Existing DL models have been trained on mainly European or East Asian populations, and the relative lack of training on darker skin pigmentation may limit overall diagnostic accuracy. This possibility is demonstrated by the increased error rates in commercial systems, trained on predominantly white datasets, for facial analysis in identifying black individuals (Buolamwini and Gebru, 2018). Second, AI may entrench existing social and economic biases and perpetuate inadvertent discriminatory practices, for example, in recommending less follow-up for black patients than for whites, when health costs are used as a proxy for health needs (Obermeyer et al., 2019). Third, disproportionate adoption by different groups may exacerbate existing inequities. Access to and use of technology differs based on sociodemographics (Tsetsi and Rains, 2017), and more techsavvy users may be more likely to embrace AI for skin screening (Tong and Sopory, 2019). The issue of equity in AI diagnosis needs to be carefully addressed to avoid inadvertent exacerbation of health-care disparities. (Young et al., 2020) - dermatology

- Model generalizability. Generalizability is a major concern for AI models; studies of computer-assisted diagnosis of melanoma report lower sensitivity for melanoma on independent test sets than on nonindependent test sets (Dick et al., 2019). It is difficult to study generalizability because published DL models are not publicly available, making it impossible to compare performance, unless each study uses a standardized benchmark database, such as the Melanoma Classification Benchmark (Brinker et al., 2019d). Han et al. (2018a) reported excellent metrics of performance and made their model available for image submission; however, the model prediction was not robust when images from an outside clinic were submitted, image magnification or contrast was altered, or images were rotated (NavarreteDechent et al., 2018). On ImageNet, a nonmedical dataset of 1,000 object categories, training on a dataset of 300 ... (Young et al., 2020)

Papers I would need access to

- https://jamanetwork.com/journals/jamadermatology/article-abstract/2688587 (linked to (Young et al., 2020))

- https://academic.oup.com/bjd/article-abstract/190/6/789/7603706?redirectedFrom=fulltext Ethical considerations for artificial intelligence in dermatology: a scoping review

2.5.4 Fairness Extensive Sources

Algorithmic Fairness

- in order to be able to fight against discrimination and achieve fairness, one should first define fairness. (Mehrabi et al., 2021)

- The fact that no universal definition of fairness exists shows the difficulty of solving this problem [138]. Different preferences and outlooks in different cultures lend a preference to different ways of looking at fairness, which makes it harder to come up with just a single definition that is acceptable to everyone in a situation. there is still no clear agreement on which constraints are the most appropriate for those problems. (Mehrabi et al., 2021)

- Broadly, fairness is the absence of any prejudice or favoritism towards an individual or a group based on their intrinsic or acquired traits in the context of decision-making [139]. Even though fairness is an incredibly desirable quality in society, it can be surprisingly difficult to achieve in practice. (Mehrabi et al., 2021)

- Here we will reiterate and provide some of the most widely used definitions, along with their explanations inspired from (Verma & Rubin, 2018).(Mehrabi et al., 2021)

- Definitions on page 12, 13, 14 (Mehrabi et al., 2021)

  - Equalized Odds (TP and FP rate should be the same for individuals in different sub groups) (Mehrabi et al., 2021)
  - Equal Opportunity (TP rate should be the same) (Mehrabi et al., 2021)
  - Demographic Parity / Statistical Parity (likelihood of positive outcome the same regardless of protected group) (Mehrabi et al., 2021)
  - Fairness Through Awareness (similar predictions to similar individuals (similarity = inverse distance)) (Mehrabi et al., 2021)
  - Fairness Through Unawareness (no protected attributes explicitly used in decision-making process) (Mehrabi et al., 2021)
  - Treatment Equality (Ration FN and FP same for both protected group categories) (Mehrabi et al., 2021)

- Test Fairness (for predicted probability scores, people in both groups must have equal probability of TP) (Mehrabi et al., 2021)
- Counterfactual Fairness (same outcome in actual world and counterfactual world where the individual belonged to a different demographic group) (Mehrabi et al., 2021)
- Fairness in Relational Domains ("A notion of fairness that is able to capture the relational structure in a domain—not only by taking attributes of individuals into consideration but by taking into account the social, organizational, and other connections between individuals" (Farnadi et al., 2018)) (Mehrabi et al., 2021) probably not relevant since not relational
- Conditional Statistical Parity (people in both groups have equal possibilities of being assigned to a positive outcome given a set of legitimate factors) (Mehrabi et al., 2021)
- Subgroup fairness: Subgroup fairness intends to obtain the best properties of the group and individual notions of fairness. It is different than these notions but uses them in order to obtain better outcomes. It picks a group fairness constraint like equalizing false positive and asks whether this constraint holds over a large collection of subgroups (Kearns et al., 2018)(Kearns et al., 2019)(Mehrabi et al., 2021)
- it is impossible to satisfy some of the fairness constraints at once except in highly constrained special cases. In [83], the authors show the inherent incompatibility of two conditions: calibration and balancing the positive and negative classes. These cannot be satisfied simultaneously with each other unless under certain constraints; therefore, it is important to take the context and application in which fairness definitions need to be used into consideration and use them accordingly [141](Mehrabi et al., 2021)
- Another important aspect to consider is time and temporal analysis of the impacts that these definitions may have on individuals or groups. In [95] authors show that current fairness definitions are not always helpful and do not promote improvement for sensitive groups—and can actually be harmful when analyzed over time in some cases. They also show that measurement errors can also act in favor of these fairness definitions; therefore, they show how temporal modeling and measurement are important in evaluation of fairness criteria and introduce a new range of trade-offs and challenges toward this direction. It is also important to pay attention to the sources of bias and their types when trying to solve fairness-related questions. (Mehrabi et al., 2021)

2.5.5 Mitigation Methods Overview

TODO: write definitions of pre-in and post-processing, see Methods for fair machine learning below [43, 11, 14]

TODO: double check and futher improve groups

Mitigation Methods - Unbiasing Data (Pre-Processing)	Mentioned in Context of	
	ML	Dermatology
Documentation and Transparency		
Good Practices while using Data	X[1,2,3]	
Datasheets as supporting document for dataset creation method, characteristics, motivations and skews	X[1,2,3]	
Datasheets as supporting document for model method, characteristics, motivations and skews	X[1,4]	
Dataset (Nutrition) Labels	X[1,5,6]	X[18, TODO: add spec source]
Communication and Reporting		
Messaging	X[1,12]	
Bias Detection and Evaluation		
Test for Simpson's Paradox TODO: Discribe Simpson's Paradox	X[1,7,8,9]	
Detect Direct Discrimination with Causal Models and Graphs	X[1,10]	
Out-of-Distribution Detection in Dermatology Using Input Perturbation and Subset Scanning		X[19]
Check confidence interval and p-curve analysis instead of p-value		X[17]
Study Design		
Allocation concealment and blinding		X[17]
Preventing Direct and Indirect Discrimination	X[1,11]	
Data Gathering		
Data Collection from diverse sources (incl. primary care clinics)	X[18]	
Robust standards for external validation	X[18]	
Preferential Sampling	X[1,13,14]	
Geographical Diversity and Inclusion for Dataset creation	X[16]	
Balanced Representation accross skin tones and genders		X[19]
Disparate Impact Removal	X[1,15]	
Labeling		
Multidimensional Scale for Skin Tones		X[19]
Data Availability and Open Science		
Publish Datasets accessible for the public		X[18, TODO: add source]

[1] (Mehrabi et al., 2021)
[2] (**M13__**)
[3] (**M55__**)
[4] (**M110__**)
[5] (**M66__**)
[6] (**M66Successor__**)
[7] (**M81__**)
[8] (**M3__**)
[9] (**M4__**)
[10] (**M163__**)
[11] (Hajian & Domingo-Ferrer, 2013)
[12] (**M74__**)
[13] (**M75__**)
[14] (**M76__**)
[15] (**M51__**)
[16] (Shankar et al., 2017)
[17] (Chakraborty, 2024)
[18] (Young et al., 2020)
[19] (Montoya et al., 2025)

Table 2.8: Mitigation Methods - Unbiasing Data - Mentioned in Contextual Research, grouped like in Mehrabi et al. (2021), the author cannot guarantee for completeness

Mitigation Methods - Fair Classification	Mentioned in Context of	
	ML	Dermatology
Satisfy Fairness Definitions		
Satisfy Subgroup Fairness TODO: unclear if * in [3] as well, or if [2] also handles *	X[1,2]	
Satisfy Equality of Opportunity*	X[1,3,6]	
Satisfy Equalized Odds*	X[1,3]	
Disparate Treatment**	X[1,4,5]	
Disparate Impact**	X[1,4,5]	
TODO: find out exact method	X[1,7]	
TODO: find out exact method	X[1,8]	
TODO: find out exact method	X[1,9]	
TODO: find out exact method	X[1,10]	
Satisfy Fairness and Stability Under Distribution Shifts		
TODO: find out exact method	X[1,11]	
Fair Representation Learning (Pre/In-processing)		
Representation Learning by Disentanglement	X[1,2]	
Variational Fair Autoencoder	X[1,3,15]	
VAE without adversarial training	X[1,4]	
Adversial Learning with FairGAN	X[1,16]	
Removing correlation between protected and unprotected features with a geometric solution	X[1,17]	
Algorithmic Adaptions for Fairness		
Modified Discrimination-Free Naive Bayes Classifier	X[1,12]	
Fairness-Aware ML Frameworks		
Fairness-Aware Classification Framework	X[1,13]	
Fairness Constraints in Multitask Learning (MTL) Framework	X[1,14]	
Decoupled Classification System with Transfer Learning	X[1,15]	
Preferential Data Selection and Representation		
Wasserstein Distance Measure for Dependence Mitigation	X[1,16]	
Preferential Sampling (PS) for Discrimination-Free Training Data	X[1,17]	
Model Interpretability		
Post-Processing with Attention Mechanism	X[1,18]	
Use Brier Score and Response Rate Accuracy		X[19], TODO: add clear source
some more methods TODO: describe		X[19]

* possible to satisfy together	[6] (**M154_**)	[13] (**M155_**)
** possible to satisfy together	[7] (**M57_**)	[14] (**M12_**)
[1] (Mehrabi et al., 2021)	[8] (**M78_**)	[15] (**M49_**)
[2] (**M147_**)	[9] (**M85_**)	[16] (**M73_**)
[3] (Hardt et al., 2016)	[10] (**M106_**)	[17] (**M75_**)
[4] (**M2_**)	[11] (**M69_**)	[18] (**M102_**)
[5] (**M159_**)	[12] (**M25_**)	[19] (Young et al., 2020)

Table 2.9: Mitigation Methods - Fair Classification - Mentioned in Contextual Research, grouped like in Mehrabi et al. (2021), the author cannot guarantee for

TODO: check categorization

Mitigation Methods - not so relevant for us	Mentioned in Context of	
	ML	Dermatology
Fair NLP		
Fair Word-Embedding	X[1,5,6,7]	
Train-Time Data Augmentation	X[1,8]	
Test-Time Neutralization	X[1,8]	
Fair Regression (In-processing)		
Price of Fairness (POF)	X[1,10]	
XY TODO: check this and bounded group loss	X[1,11]	
Decision Tree for Disparate Impact and Treatment	X[1,12]	
Structured Prediction (In-processing)		
Reducing Bias Amplification (RBA) as calibration algorithm	X[1,13]	
Principal Component Analysis (PCA) (In-processing)		
Fair PCA	X[1,14]	
Graph-Based Fairness Methods		
Community Detection / Graph Embedding TODO: how to proceed with this	X	
Causal Fairness and Disparate Learning		
Disparate Learning Processes (DLP)	X[1,9]	
Causal Approach to Fairness TODO: how to proceed with this	X[TODO: add clear source]	
Disregard path in causal graph which result in sensitive attributes affecting decision outcome	X[1]	
Removing Sensitive Attributes		
Disregard sensitive attributes in effect on decision making	X[1]	

[1] (Mehrabi et al., 2021)
[2] (**M42_**)
[3] (**M97_**)
[4] (**M112_**)
[5] (Bolukbasi et al., 2016)
[6] (**M58_**)
[7] (**M169_**)
[8] (**M166_**)
[9] (**M94_**)
[10] (**M14_**)
[11] (**M1_**)
[12] (**M2_**)
[13] (Zhao et al., 2017)
[14] (**M137_**)
[15] (**M5_**)
[16] (**M90_**)
[17] (**M65_**)

Table 2.10: Mitigation Methods - Others - Mentioned in Contextual Research, grouped like in Mehrabi et al. (2021), the author cannot guarantee for completeness

TODO: mention also the IBM AI Fairness 360 toolkit [11] and that authors evaluated their work in benchmark datasets [65], [72], [158], [159]

TODO: draft for presentation satisfy Equalized Odds / Subgroup fairness highlight allocation concealment and blinding and data collection from diverse sources

and Preferential Sampling

2.5.6 Mitigation Methods Overview

Mitigation Methods	Mentioned in Context of	
	ML	Dermatology
Unbiasing Data		
Documentation and Transparency	X[1]	X[3]
Bias Detection and Evaluation	X[1]	X[2,4]
Study Design	X[1]	X[2]
Data Gathering	X[1]	X[3,4]
Data Availability and Open Science		X[3]
Removing Sensitive Attributes	X[1]	
Fair Classification		
Satisfy Fairness Definitions	X[1]	
Satisfy Fairness and Stability Under Distribution Shifts	X[1]	
Fair Representation Learning	X[1]	
Fairness-Aware ML Frameworks	X[1]	
Preferential Data Selection and Representation	X[1]	
Model Interpretability	X[1]	X[3]
For Other ML Algorithm Types		
Fair NLP	X[1]	
Fair Regression	X[1]	
Structured Prediction	X[1]	
Fair Principal Component Analysis	X[1]	
Graph-Based Fairness Methods	X[1]	
Causal Fairness and Disparate Learning	X[1]	

[1] (Mehrabi et al., 2021)   [3] (Young et al., 2020)   [4] (Montoya et al., 2025)
[2] (Chakraborty, 2024)

Table 2.11: Mitigation Methods - Draft

TODO: check categorization

Mitigation Methods - not so relevant for us	Mentioned in Context of	
	ML	Dermatology
Fair NLP		
Fair Word-Embedding	X[1,5,6,7]	
Train-Time Data Augmentation	X[1,8]	
Test-Time Neutralization	X[1,8]	
Fair Regression (In-processing)		
Price of Fairness (POF)	X[1,10]	
XY TODO: check this and bounded group loss	X[1,11]	
Decision Tree for Disparate Impact and Treatment	X[1,12]	
Structured Prediction (In-processing)		
Reducing Bias Amplification (RBA) as calibration algorithm	X[1,13]	
Principal Component Analysis (PCA) (In-processing)		
Fair PCA	X[1,14]	
Graph-Based Fairness Methods		
Community Detection / Graph Embedding TODO: how to proceed with this	X	
Causal Fairness and Disparate Learning		
Disparate Learning Processes (DLP)	X[1,9]	
Causal Approach to Fairness TODO: how to proceed with this	X[TODO: add clear source]	
Disregard path in causal graph which result in sensitive attributes affecting decision outcome	X[1]	
Removing Sensitive Attributes		
Disregard sensitive attributes in effect on decision making	X[1]	

[1] (Mehrabi et al., 2021) [7] (**M169__**) [13] (Zhao et al., 2017)
[2] (**M42__**) [8] (**M166__**) [14] (**M137__**)
[3] (**M97__**) [9] (**M94__**) [15] (**M5__**)
[4] (**M112__**) [10] (**M14__**) [16] (**M90__**)
[5] (Bolukbasi et al., 2016) [11] (**M1__**) [17] (**M65__**)
[6] (**M58__**) [12] (**M2__**)

Table 2.12: Mitigation Methods - Others - Mentioned in Contextual Research, grouped like in Mehrabi et al. (2021), the author cannot guarantee for completeness

2.5.7   Mitigation Methods Extensive Sources

Bias Examples and Mitigation Ideas

Data bias examples and mitigation ideas

- Bias in ML Data - (Buolamwini & Gebru, 2018) IJB-A / Adience imbalanced (mainly light-skinned subjects) - Bias towards dark-skinned groups (under-represented). Other instance - when we do not consider different subgroups in the data. Considering only male-female groups not enough, use race to further subdivide gender groups. Only then, clear biases in sub groups can be found, since otherwise part of the groups would compromise the other group and hide the underlaying bias towards that subgroup (Mehrabi et al., 2021)

- Popular machine-learning datasets that serve as a base for most of the developed algorithms and tools can also be biased—which can be harmful to the downstream applications that are based on these datasets. ... In [(Shankar et al., 2017), researchers showed that these datasets suffer from representation bias and advocate for the need to incorporate geographic diversity and inclusion while creating such datasets. (Mehrabi et al., 2021)

- Examples of Data Bias in Medical Applications. These data biases can be more dangerous in other sensitive applications. For example, in medical domains there are many instances in which the data studied and used are skewed toward certain populations—which can have dangerous consequences for the underrepresented communities. [98] showed how exclusion of African-Americans resulted in their misclassification in clinical studies, so they became advocates for sequencing the genomes of diverse populations in the data to prevent harm to underrepresented populations (Mehrabi et al., 2021) TODO: What does sequencing data mean?, is it relevant

Methods for Fair Machine Learning

- While this section is largely domain-specific, it can be useful to take a cross-domain view. Generally, methods that target biases in the algorithms fall under three categories (Mehrabi et al., 2021)

- Pre-processing. Pre-processing techniques try to transform the data so that the underlying discrimination is removed [43]. If the algorithm is allowed to modify the training data, then pre-processing can be used [11].(Mehrabi et al., 2021)

- In-processing. In-processing techniques try to modify and change state-of-the-art learning algorithms in order to remove discrimination during the model training process [43]. If it is allowed to change the learning procedure for a machine learning model, then in-processing can be used during the training of a model— either by incorporating changes into the objective function or imposing a constraint [11, 14].(Mehrabi et al., 2021)

- Post-processing. Post-processing is performed after training by accessing a holdout set which was not involved during the training of the model [43]. If the algorithm can only treat the learned model as a black box without any

ability to modify the training data or learning algorithm, then only post-processing can be used in which the labels assigned by the black-box model initially get reassigned based on a function during the post-processing phase [11, 14].(Mehrabi et al., 2021)

- we concentrate on discrimination prevention based on preprocessing, because the preprocessing approach seems the most flexible one: it does not require changing the standard data mining algorithms, unlike the inprocessing approach, and it allows data publishing (rather than just knowledge publishing), unlike the postprocessing approach. (Hajian & Domingo-Ferrer, 2013) –> TODO: this is an important point which we should consider for PASSION, also, some more insight in regards of the different phases can be found in this paper

- From learning fair representations [42, 97, 112] to learning fair word embeddings [(Bolukbasi et al., 2016), 58, 169], debiasing methods have been proposed in different AI applications and domains. (Mehrabi et al., 2021) –> seems to refer mostly to NLP domains

- Most of these methods try to avoid unethical interference of sensitive or protected attributes into the decision-making process, while others target exclusion bias by trying to include users from sensitive groups. (Mehrabi et al., 2021)

- However, a recent paper [58] argues against these debiasing techniques and states that many recent works on debiasing word embeddings have been superficial, that those techniques just hide the bias and don't actually remove it. (Mehrabi et al., 2021)

- some works try to satisfy one or more of the fairness notions in their methods, such as disparate learning processes (DLPs) which try to satisfy notions of treatment disparity and impact disparity by allowing the protected attributes during the training phase but avoiding them during prediction time [94].(Mehrabi et al., 2021)

- Some of the existing work tries to treat sensitive attributes as noise to disregard their effect on decision-making, while some causal methods use causal graphs, and disregard some paths in the causal graph that result in sensitive attributes affecting the outcome of the decision.(Mehrabi et al., 2021)

- Different bias-mitigating methods and techniques are discussed below for different domains—each targeting a different problem in different areas of machine learning in detail. (Mehrabi et al., 2021)

Unbiasing Data

- Every dataset is the result of several design decisions made by the data curator. Those decisions have consequences for the fairness of the resulting

dataset, which in turn affects the resulting algorithms. In order to mitigate the effects of bias in data, some general methods have been proposed that advocate having good practices while using data, such as having datasheets that would act like a supporting document for the data reporting the dataset creation method, its characteristics, motivations, and its skews [13, 55]. A similar suggestion has been proposed for models in [110].(Mehrabi et al., 2021)

- Authors in [66] also propose having labels, just like nutrition labels on food, in order to better categorize each data for each task. (Mehrabi et al., 2021)

- some work has targeted more specific types of biases. For example, [81] has proposed methods to test for cases of Simpson's paradox in the data, and [3, 4] proposed methods to discover Simpson's paradoxes in data automatically. (Mehrabi et al., 2021)

- Causal models and graphs were also used in some work to detect direct discrimination in the data along with its prevention technique that modifies the data such that the predictions would be absent from direct discrimination [163].(Mehrabi et al., 2021)

- in [(Hajian & Domingo-Ferrer, 2013)] also worked on preventing discrimination in data mining, targeting direct, indirect, and simultaneous effects.(Mehrabi et al., 2021)

- Other pre-processing approaches, such as messaging [74], preferential sampling [75, 76], disparate impact removal [51], also aim to remove biases from the data. (Mehrabi et al., 2021)

- Image quality. Several barriers to AI implementation in the clinic need to be overcome with regards to imaging (Figure 1). These include technical variations (e.g., camera hardware and software) and differences in image acquisition and quality (e.g., zoom level, focus, lighting, and presence of hair). For example, the presence of surgical ink markings is associated with decreased specificity (Winkler et al., 2019), field of view can significantly affect prediction quality (Mishra et al., 2019), and classification performance improves when hair and rulers are removed (Bisla et al., 2019). We have developed a method to measure how model predictions might be biased by the presence of a visual artifact (e.g., ink) and proposed methods to reduce such biases (Pfau et al., 2019). Poor quality images are often excluded from studies, but the problem of what makes an image adequate is not well studied. Ideally, models need to be able to express a level of confidence in a prediction as a function of image quality and appropriately direct a user to retake photos if needed. (Young et al., 2020) - dermatology

Fair Classification

- certain methods have been proposed [57, 78, 85, 106] that satisfy certain definitions of fairness in classification. For instance, in [147] authors try to satisfy subgroup fairness in classification, equality of opportunity and equalized odds in [63], both disparate treatment and disparate impact in [2, 159], and equalized odds in [154]. (Mehrabi et al., 2021)

- Other methods try to not only satisfy some fairness constraints but to also be stable toward change in the test set [69] (Mehrabi et al., 2021)

- The authors in [155], propose a general framework for learning fair classifiers. This framework can be used for formulating fairness-aware classification with fairness guarantees. In another work [25], authors propose three different modifications to the existing Naive Bayes classifier for discrimination-free classification.(Mehrabi et al., 2021)

- paper [122] takes a new approach into fair classification by imposing fairness constraints into a Multitask learning (MTL) framework. In addition to imposing fairness during training, this approach can benefit the minority groups by focusing on maximizing the average accuracy of each group as opposed to maximizing the accuracy as a whole without attention to accuracy across different groups. In a similar work [49], authors propose a decoupled classification system where a separate classifier is learned for each group. They use transfer learning to reduce the issue of having less data for minority groups.(Mehrabi et al., 2021)

- In [73] authors propose to achieve fair classification by mitigating the dependence of the classification outcome on the sensitive attributes by utilizing the Wasserstein distance measure.(Mehrabi et al., 2021)

- In [75] authors propose the Preferential Sampling (PS) method to create a discrimination free train data set. They then learn a classifier on this discrimination free dataset to have a classifier with no discrimination.(Mehrabi et al., 2021)

- In [102], authors propose a post-processing bias mitigation strategy that utilizes attention mechanism for classification and that can provide interpretability. (Mehrabi et al., 2021)

Fair Regression   TODO: only summarize briefly, as PASSION is a classification and not a regression task

- "price of fairness" (POF) to measure accuracy-fairness trade-offs, 3 penalites: Individual fairness, group fairness and hybrid fairness [14] (Mehrabi et al., 2021)

- In addition to the previous work, [1] considers the fair regression problem formulation with regards to two notions of fairness statistical (demographic) parity and bounded group loss. [2] uses decision trees to satisfy disparate impact and treatment in regression tasks in addition to classification. (Mehrabi et al., 2021)

Structured Prediction   TODO: only summarize briefly, as PASSION is a classification task

- RBA (reducing bias amplification) as calibration algorithm to prevent risk of leveraging social bias, distributions in training data are followed in the predictions. multi-label obeject and visual semantic role labeling classification amplify existing bias in data [(Zhao et al., 2017)] (Mehrabi et al., 2021) –> TODO: be careful with this if the approach would be to generate new images for training!!

Fair PCA   TODO: only summarize briefly, as PASSION is a classification task with only like 10 features

- Pincipal Component Analysis (PCA) https://www.geeksforgeeks.org/principal-component-analysis-pca/ –> dimensionality reduction, statistical technic, high-dimensional data into lower-dimensional space while maximising variance in new space -> most important patterns and relationships is preserved

- vanilla PCA exaggerate error in reconstruction for one group of people [137] (Mehrabi et al., 2021)

- And their proposed algorithm is a two-step process listed below: (1) Relax the Fair PCA objective to a semidefinite program (SDP) and solve it. (2) Solve a linear program that would reduce the rank of the solution. [137] (Mehrabi et al., 2021)

Community Detection   TODO: use this as an example for out of scope text, - Ludovic approved Community detection algorithms are specifically tailored to analyze network data and find connections in such datasets. For example, they can be used to detect groups of people with similar interest in social networks (Jayawickrama, 2021). This kind of data is not found in the context of PASSION, which is a classification task. Please refer to Mehrabi et al. (2021) for more information on bias mitigation in community detection algorithms.

Causal Approach to Fairness   TODO: only relevant, if our variables have a dependency on the variables, e.g. age / gender determines how the disease is presenting itself in the images; check (Mehrabi et al., 2021) page 18 if relevant

Fair Representation Learning   https://medium.com/superlinear-eu-blog/representation-learning-breakthroughs-what-is-representation-learning-5dda2e2fed2e

- Variational Auto encoders –> Variational Fair Autoencoder introduced in [97]. Here,they treat the sensitive variable as the nuisance variable, so that by removing the information about this variable they will get a fair representation. They use a maximum mean discrepancy regularizer to obtain invariance in the posterior distribution over latent variables. Adding this maximum mean discrepancy (MMD) penalty into the lower bound of their

VAE architecture satisfies their proposed model for having the Variational Fair Autoencoder.

In [5] authors also propose a debiased VAE architecture called DB-VAE which learns sensitive latent variables that can bias the model (e.g., skin tone, gender, etc.) and propose an algorithm on top of this DB-VAE using these latent variables to debias systems like facial detection systems.

In [112] authors model their representation-learning task as an optimization objective that would minimize the loss of the mutual information between the encoding and the sensitive variable. The relaxed version of this assumption is shown in Equation 1. They use this in order to learn fair representation and show that adversarial training is unnecessary and in some cases even counter-productive.

In [42], authors introduce flexibly fair representation learning by disentanglement that disentangles information from multiple sensitive attributes. Their flexible and fair variational autoencoder is not only flexible with respect to downstream task labels but also flexible with respect to sensitive attributes. They address the demographic parity notion of fairness, which can target multiple sensitive attributes or any subset combination of them. (Mehrabi et al., 2021)

- Adversarial Learning - In [90] authors present a framework to mitigate bias in models learned from data with stereotypical associations. using adversarial networks by introducing FairGAN which generates synthetic data that is free from discrimination and is similar to the real data. They use their newly generated synthetic data from FairGAN, which is now debiased, instead of the real data for training and testing. They do not try to remove discrimination from the dataset, unlike many of the existing approaches, but instead generate new datasets similar to the real one which is debiased and preserves good data utility. (Mehrabi et al., 2021) TODO: address challenges in creating synthetic data in dermatology?

Fair NLP   TODO: for PASSION irrelevant, if it wants to stick to ResNet50 Architecture (Gottfrois et al., 2024) and not use Visual Encoders, which would make sense bc of the small dataset

- Word Embedding TODO: potentially relevant, if the labels are used in training, e.g. age / gender determines how the disease is presenting itself in the images; check (Mehrabi et al., 2021) page 21 if relevant

- Coreference Resolution "Coreference resolution involves identifying when two or more expressions in a text refer to the same entity, be it a person, place, or thing." https://medium.com/@datailm/the-key-to-unlocking-true-language-understanding-coreference-resolution-c01d569e2e87 TODO: irrelevant for the PASSION Context

comparison of different mitigation algorithms

- The field of algorithmic fairness is a relatively new area of research and work still needs to be done for its improvement. With that being said, there are already papers that propose fair AI algorithms and bias mitigation techniques and compare different mitigation algorithms using different benchmark datasets in the fairness domain. For instance, authors in [65] propose a geometric solution to learn fair representations that removes correlation between protected and unprotected features. The proposed approach can control the trade-off between fairness and accuracy via an adjustable parameter. In this work, authors evaluate the performance of their approach on different benchmark datasets, such as COMPAS, Adult and German, and compare them against various different approaches for fair learning algorithms considering fairness and accuracy measures [65, 72, 158, 159]. In addition, IBM's AI Fairness 360 (AIF360) toolkit [11] has implemented many of the current fair learning algorithms and has demonstrated some of the results as demos which can be utilized by interested users to compare different methods with regards to different fairness measures. (Mehrabi et al., 2021)

2.5.8 Statistical biases

https://data36.com/statistical-bias-types-explained/

-

2.5.9 Dermatology Bias

- https://ijdvl.com/biases-in-dermatology-a-primer/ 29 biases, 4 reasons to know about it, 7 mitigation methods (Chakraborty, 2024) - dermatology

- A recent study reported mean top-1 and top-5 model accuracy of 44.8% and 78.1%, respectively, for the classification of 134 diseases (Han et al., 2019b). Most datasets are proprietary, often with minimal description, and datasets collected in dermatology clinics may be skewed toward more complex cases, to those patients with better access to care, or by the choice of camera used in one clinic versus another. Data should be collected from as many diverse sources as possible, including primary care clinics, and robust standards for external validation are needed. (Young et al., 2020)

- There have been successful efforts to support reproducibility and open access. For example, the study by Han et al. (2018a) details the number and characteristics of images from each data source and makes thumbnails of the images publicly available. Additionally, several studies classifying dermoscopic images use the publicly available International Skin Imaging Collaboration archive (Gutman et al., 2016). By making datasets public, it becomes possible to examine them for bias (Bissoto et al., 2019). Alternatively, reporting a model training database's patient demographics and

disease classes would be helpful in predicting model performance on external populations. (Young et al., 2020)

- Metrics of model performance. Standard metrics are needed to assess the performance of different models (Figure 1). Currently, standard performance metrics such as accuracy and area under the receiver operating characteristic and precision recall curves are routinely reported. However, for use in the clinic, studies should additionally describe how well their models deal with uncertainty by reporting (i) the Brier Score, or mean-squared calibration error (Rufibach, 2010), which measures how reliably a model can forecast its accuracy, and (ii) area under the response rate accuracy curve, which measures how capably a model can identify examples it is likely to predict falsely and thus abstain from predicting (Hendrycks et al., 2019) (Young et al., 2020)

- Model interpretability. Acceptance of AI in clinical decision making hinges on being able to understand the decisionmaking process fundamental to its predictions. DL models are inherently difficult to interpret because they are complex, routinely containing millions of learned parameters; interpretation of DL models' output is an active field of research (Murdoch et al., 2019). One approach for interpreting model diagnoses is contentbased image retrieval, a method for retrieving training images that are visually similar to a test image (Tschandl et al., 2019a). This method may reassure the physician if all the retrieved training images have the same diagnosis as the predicted diagnosis but is less helpful if the test image looks similar to two or more training images with conflicting diagnoses. A second approach is to highlight pixels in an image most relevant for a model's prediction, using methods such as saliency mapping (Figure 1). However, it is often the case that highlighted pixels correspond to the entire lesion or visually distinctive features that are already obvious to clinicians without indication as to why these pixels are important to the diagnosis. A third approach is to see through the eyes of a model by plotting an activation atlas (Carter et al., 2019), which shows how subtle changes, in particular visual features, may tip the model over into choosing one diagnosis over another. These activation atlases are experimental and have yet to be applied in dermatology. Understanding a model's predictions and how the prediction is applicable to the patient at hand is necessary to build trust. As AI exceeds human performance in various tasks, interpreting models may help to advance scientific knowledge by understanding what the machine sees that is relevant to its predications (Young et al., 2020)

2.5.9.1 Demographic Bias in Dermatology

fairness melanoma detection

- Some biases can be easily detected and countered, such as through appropriate data curation; for example, having a balanced representation across

skin tones and genders in training sets. However, in other cases, biases are hidden and untraceable [9]. (Montoya et al., 2025)

- whether information on demographic diversity (age, gender, race, or ethnicity of patients), clinical diversity (skin type, lesion type, anatomical location of lesion), or image characteristics (source, imaging techniques, resolution, and whether the images were real or artificially generated) (Montoya et al., 2025)

- The most popular skin color scale currently being used for data annotation for image recognition techniques is the Fitzpatrick Skin Tone Scale (FST) [10]which has six skin tones. Dating from the 1970s, it originally featured just 4 light tones and was designed for detecting photo sensitivity for white skin, with two darker tones added later [11]. The Monk Skin Scale was recently developed and still needs testing, but promisingly has 10 tones, 5 light and 5 dark [12].(Montoya et al., 2025) TODO: highlight this (FST alternatives)

- Fig. 4. Comparison of skin tone scales that can be used for skin cancer detection utilizing AI. Recreation of fitzpatrick skin type scale, monk skin tone scale, and sampling of L'Oreal color chart map for reference. (Montoya et al., 2025) TODO: include this figure

- While this systemic review provides a comprehensive review of the literature on fairness in AI for melanoma detection, it is primarily based on existing research. To validate the proposed recommendations or frameworks, continuing work is necessary to complete empirical analysis and experiments. Additionally, the suggested adoption of new skin tone scales, while beneficial, may face practical challenges in implementation. Furthermore, while the paper strongly advocates for specific skin tone scales, it's important to note that other methods or tools might also effectively address fairness issues in AI for melanoma detection. Finally, while the study addresses fairness in AI, it could benefit from further exploration of the practical implementation of these recommendations in real-world clinical settings. Potential obstacles and the feasibility of widespread adoption should be considered to ensure that the proposed solutions are not only theoretically sound but also practically viable. (Montoya et al., 2025) TODO: also mention the limitations regarding FST alternatives

- Recent research [13] adds another axis, skin hue, which is described as ranging from red to yellow. This offers a more complete representation of variations of skin color by providing a multidimensional scale [13]. (Montoya et al., 2025)

- The effect of hue (blue, red, yellow, green) on skin tones adds depth to each face producing a range of undertones (cold, neutral, warm, and olive). In the realm of color theory, the concept of 'contrast of hue' emphasizes the distinctiveness among fundamental colors, with primary hues like yellow, red, and blue exhibiting the most pronounced differences [14]. Because skin cancer

appears differently on different colored skin, it is important to acknowledge a full range of colors present in both healthy skin and suspicious lesions within datasets used to train skin cancer detection ML tools. (Montoya et al., 2025)

- These findings should correlate to AI for melanoma detection since the contrast between skin color and skin lesions is a preliminary marker during feature extraction. Although the Fitzpatrick Skin Tone (FST) FST measurement scale is not diverse enough and leads to biased AI tools, it is continually used and has even been used to test a recently FDA-approved AI device for detecting melanoma. (Montoya et al., 2025)

- We advocate for the adoption of improved scales like the Monk and L'Oreal maps. Future studies should ensure equitable representation and testing across skin tones to guarantee AI's effectiveness for all. Please refer to Tables 2 through 7 in the discussion section for further recommendations for curating a diverse dataset, including purpose, ownership, funding, and data annotation, as well as recommendations for each stage of the data life cycle. (Montoya et al., 2025) TODO: Link for further mitigation methods

- This study found that while using skin tone instead of race for fairness evaluations in computer vision seems objective, the annotation process remains biased by human annotators. Untested scales, unclear procedures, and a lack of awareness about annotator backgrounds and social context significantly influence skin tone labeling. This study exposes how even minor design choices in the annotation process, like scale order (dark to light instead of light to dark) or image context (face or no face, skin lesion presence), can sway agreement and introduce uncertainty in skin tone assessments. ... The researchers emphasize the need for greater transparency, standardized procedures, and careful consideration of annotator biases to mitigate these challenges and ensure fairer and more robust evaluations in computer vision. (Montoya et al., 2025) - demographic dermatology bias

3 Ideas and Concepts

Hier geht es um die Fragestellung, wie Sie die formulierten Ziele der Arbeit erreichen wollen. Sie halten z.B. erste, grobe Ideen, skizzenhafte Lösungsansätze fest. Gibt es mehrere Wege, Ansätze um dieses Ziel zu erreichen, begründen Sie hier, warum Sie einen bestimmten Weg einschlagen. Beispiel für ein Softwareprojekt: Erste Gedanken über eine grobe Systemarchitektur. Ist z.B. eine Microservice-Architektur angebracht? Welche Alternativen bestehen, wo gibt es Problempunkte? Die Umsetzung, die Beurteilung der Machbarkeit und die detaillierte Beschreibung der umgesetzten Architektur sind dann Teil der Realisierung.

3.1 PASSION Dataset

TODO: write things to consider more precisely:

- Include more details in gender attribute - transgender have probably different genes / hormones, and should be indicated for more accurraccy

- include profession / at least an adapted version to indicate high risk patients for certain diseases? -> might lead to other biases?

- change country of origin to ethnicity (less of a proxy variable)

- are the data collectors specialized in some fields? That could lead to bias towards the center's country and the diagnosed diseases

- include images of healthy skin

3.2 Broad Methodology

TODO: should this really be stated here or in the methodology section? First, I need to gain an overview over what biases, fairness metrics and mitigation strategies are known in general. Then, I must scope the found information, to find what is relevant for PASSION and what is feasible to achieve within the time constraints of this thesis. Before starting an assessment, a baseline needs to be established by computing chosen fairness metrics. Afterwards, the mitigation methods can be applied, and the performance can be compared to the baseline, to find out whether the methods are indeed mitigating the biases.

TODO: add infos from the midterm presentation
TODO: write things to consider more precisely:

- Divide and Conquer vs. All-In-One-Model

  - An algorithm per ethnicity / subgroup running at the same time

  - Running 1 Algorithm chosen based on Fitzpatrick skin type

  - Running 1 Algorithm which detects first the demographic subgroup (FST, gender, age, …) and runs the specific subgroup algorithm afterwards

  - Hint Ludovic: Still not of data, maybe also others; often limited because the data is missing, you are missing data from others

- BLIND performance vs. Including the demographic data

  - Idea to try if the labels are not relevant for the diagnosis and should only be used for evaluating fairness purposes as some papers suggest

  - Might be obsolete after demographic biases in dermatology research, since melanin response and melanoma risk is different in male and female according to research https://pmc.ncbi.nlm.nih.gov/articles/PMC4797181/

- Hint Ludovic: Maybe Focal Loss more relevant –> emphasis on data vs model

- Divide and Conquer vs. All-In-One-Model (either by etnicity x algorithms at a time or one which seperates the imgs first by demographic subgroup (incl. Fitzpatrick skin type))

- BLIND performance vs. Including the demographic data

4 Methods

Hier halten Sie fest und begründen, welches Vorgehensmodell Sie für Ihr Projekt wählen. Sie verweisen allenfalls auf die daraus entstandenen, konkreten Terminpläne mit Meilensteinen, welche z.B. unter Realisierung (Kapitel 5) oder im Anhang versorgt sind. Bei Projekten mit einer verlangten wissenschaftlichen Tiefe werden hier die geplanten Forschungsmethoden wie quantitative/qualitative Interviews, Befragungen, Beobachtungen, Feldexperiment etc. beschrieben und begründet. Warum ist in Ihrer Situation ein Interview besser als eine Umfrage? Wer soll interview werden? Die gewählten Methoden sind nachvollziehbar und begründet. Eine methodische Übersicht (Methodisches BigPicture) wurde aufgezeigt und Abgrenzungen erläutert.

4.0.1 Bias Evaluation in PASSION

TODO: fix this writing

Methodology

In order to evaluate which biases are there in PASSION dataset, I need to

- reproduce the results of the paper using the PASSION evaluation project with the PASSION data to see a) verify the paper results, b) check what analysis data is available for the evaluation (the paper provides probably a summary), c) check what code can be reused and what needs to be adapted

- evaluate what data is missing to do evaluate the fairness and biases

- adapt code so that the relevant data is generated to be able to compute the relevant information

- GPUHub from HSLU is used, TODO: maybe add some machine information

Execution

In order to run the reproduction of the PASSION results on GPUHub, some minor changes in the loading process for the metadata had to be done. The required changes will be contributed to the code base to make the reproduction easier for others. TODO: do that then ;) All 4 experiments where ran to get a broad overview over the results. The results of the 4 experiments contained the same scores as the ones mentioned in the paper. The scores are given separately for each

demographic groups, either based on skin type or gender. There is no data for subgroups available. TODO: add results from age and center experiments For the conditions classification, the provided scores per demographic group in the paper are the average scores over all classes. The variance/deviation of the averages in the scores based on the condition varies between the groups. E.g. the F1 score for FST VI is 0.71 ± 0.11 (total support 87) while for FST it is 0.73 ± 0.04 (total support 254). The detailed information can be found in the attached logs TODO: add logs from C:

Users

nadja

OneDrive

HSLU_Nadja

BAA

baa_on_git

results

reproducing_PASSION_results.

The available data allows to compute the equalized odds fairness. TODO: describe which fairness methods to use why In order to be able compute the subgroup fairness, the test results need to be split further into the subgroups. TODO: check exactly how to calculate subgroup fairness and whether there is already an algorithm for it, same for equalized odds

To reproduce all 4 experiments, the training ran took roughly TODO: add: it started on 24.4. roughly at 7 o clock, took until ca. 16.4., 13:00, the two other experiments started on 26.4. 17:30 hours since there where no model checkpoints available. This runtime is not feasible for each new training using a mitigation method, especially because the GPUHub seems to cancel the script after roughly 2 days, which means that each experiment needs to be started independently. Since this runtime it is not feasible to use for every mitigation run, the following adaptations where made TODO: add improvements described in protokol week 9

checkpoint handling was a bit broken runs x and x1 showed, that my code was reproducible while theirs was wrong –> running on passion model checkpoint without loading the rest of the infos in the checkpoint —-> confirmed, that my calculations are reproducible while theirs are not; see commented out file names for comparison also tested on another checkpoint, same behaviour, not cached results though

no implementation of subgroup fairness found –> calc eq odds for subgroups

The results in the PASSION results where reproduced by After adapting the data linkage in the PASSION evaluation code, the results where available whe in the PASSION evaluation code.

Evaluation and Validation

5 Execution

Dies ist das Hauptkapitel Ihrer Arbeit! Hier wird die Umsetzung der eigenen Ideen und Konzepte (Kapitel 3) anhand der gewählten Methoden (Kapitel 4) beschrieben, inkl. der dabei aufgetretenen Schwierigkeiten und Einschränkungen. Die gewählten Methoden werden systematisch, konsistent und korrekt auf den Kontext der Arbeit angewendet. Die Bearbeitungs- bzw. Forschungsobjekte sind einheitlich benannt, im Kontext dargestellt und sinnvoll in die Arbeit integriert. Praxis- und Erfahrungswissen (z.B. aus Interviews) wird zur Validierung und Ergänzung der erarbeiteten Ergebnisse herangezogen.

6 Evaluation and Validation

Auswertung und Interpretation der Ergebnisse. Nachweis, dass die Ziele erreicht wurden, oder warum welche nicht erreicht wurden. Die Ziele / Forschungsfragen sind dem Umfang der Arbeit entsprechend sehr klar abgegrenzt; sie sind präzise, überprüfbar und nach den Standards der Zielformulierung definiert. Die Zielerreichung wurde systematisch und korrekt validiert. Die Herleitung und Bedeutung der Ergebnisse, mögliche Varianten, Gütekriterien und eine Validierung allgemein werden nachvollziehbar diskutiert

7 Outlook

Reflexion der eigenen Arbeit, ungelöste Probleme, weitere Ideen. Die Ergebnisse und Empfehlungen schaffen einen konkreten Mehrwert für die Auftraggebenden. Einschränkungen und Grenzen werden kritisch diskutiert und die nächsten Schritte im Ausblick festgehalten, so dass die Ergebnisse direkt in der Praxis weiterverwendet und/oder angewendet werden können.

8 GlossaryTODO: probably remove this

TODO: Add List of Formulas if necessary TODO: add AI declarations somewhere

9 Bibliography

Aleid, A. M., Nukaly, H. Y., Almunahi, L. K., Albwah, A. A., AL-Balawi, R. M. D., AlRashdi, M. H., Alkhars, O. A., Alrasheeday, A. M., Alshammari, B., Alabbasi, Y., & Al Mutair, A. (2024). Prevalence and socio-demographic and hygiene factors influencing impetigo in saudi arabian children: A cross-sectional investigation [Publisher: Dove Medical Press eprint: https://www.tandfonline.com *Clinical, Cosmetic and Investigational Dermatology*, *17*, 2635–2648. https://doi.org/10.2147/CCID.S472228

Baeza-Yates, R. (2018). Bias on the web. *Commun. ACM*, *61*(6), 54–61. https://doi.org/10.1145/3209581
Mehrabi 9.

Baldé, B. (2023, April 14). *Why you should use stratified split* [Medium]. Retrieved April 14, 2025, from https://medium.com/@becaye-balde/why-you-should-use-stratified-split-bddb6dadd34e

Berk, R., Heidari, H., Jabbari, S., Kearns, M., & Roth, A. (2017). Fairness in criminal justice risk assessments: The state of the art [Publisher: SAGE Publications Inc]. *Sociological Methods & Research*, *50*(1), 3–44. https://doi.org/10.1177/0049124118782533
Mehrabi 15.

Bolukbasi, T., Chang, K.-W., Zou, J. Y., Saligrama, V., & Kalai, A. T. (2016). Man is to computer programmer as woman is to homemaker? debiasing word embeddings. *Advances in Neural Information Processing Systems*, *29*. Retrieved April 3, 2025, from https://proceedings.neurips.cc/paper_files/paper/2016/hash/a486cd07e4ac3d270571622f4f316ec5-Abstract.html
Mehrabi 20.

British Association of Dermatologists (BAD). (2021, July 7). *Lower socioeconomic status linked with more severe skin disease, including melanoma* [Bad patient hub] [Research was presented at the BAD's Annual Meeting.]. Retrieved February 17, 2025, from https://www.skinhealthinfo.org.uk/lower-socioeconomic-status-linked-with-more-severe-skin-disease-including-melanoma/

Buolamwini, J., & Gebru, T. (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification [ISSN: 2640-3498]. *Proceedings of the 1st Conference on Fairness, Accountability and Transparency*, 77–91. Retrieved March 16, 2025, from https://proceedings.mlr.press/v81/buolamwini18a.html
Mehrabi 24, demographic (skin type and gender).

Chakraborty, A. (2024). Biases in dermatology: A primer [Publisher: Scientific Scholar]. *Indian J Dermatol Venereol Leprol, 90*(2), 250–254. https://doi.org/10.25259/IJDVL_126_2023
0 citations (but from 2024), list of lots of biases.

Chen, J., Kallus, N., Mao, X., Svacha, G., & Udell, M. (2019). Fairness under unawareness: Assessing disparity when protected class is unobserved. *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 339–348. https://doi.org/10.1145/3287560.3287594
Mehrabi 30.

Chouldechova, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments [Publisher: Mary Ann Liebert, Inc., publishers]. *Big Data, 5*(2), 153–163. https://doi.org/10.1089/big.2016.0047
Mehrabi 34.

Ciampaglia, G. L., Nematzadeh, A., Menczer, F., & Flammini, A. (2018). How algorithmic popularity bias hinders or promotes quality [Publisher: Nature Publishing Group]. *Sci Rep, 8*(1), 15951. https://doi.org/10.1038/s41598-018-34203-2
Mehrabi 117.

Clarke, K. A. (2005). The phantom menace: Omitted variable bias in econometric research [Publisher: SAGE Publications Ltd]. *Conflict Management and Peace Science, 22*(4), 341–352. https://doi.org/10.1080/07388940500339183
Mehrabi 38, difficultis regarding ommitted variable and overcoming methods.

Corbett-Davies, S., Pierson, E., Feller, A., Goel, S., & Huq, A. (2017). Algorithmic decision making and the cost of fairness. *Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 797–806. https://doi.org/10.1145/3097983.3098095
Mehrabi 41.

Danks, D., & London, A. J. (2017). Algorithmic bias in autonomous systems. *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence*, 4691–4697. https://doi.org/10.24963/ijcai.2017/654
Mehrabi 44.

Delgado-Rodríguez, M., & Llorca, J. (2004). Bias [Publisher: BMJ Publishing Group Ltd Section: Continuing professional education]. *Journal of Epidemiology & Community Health, 58*(8), 635–641. https://doi.org/10.1136/jech.2003.008466

Diaz, M., Lucke-Wold, B., Batchu, S., & Kleinberg, G. (2022). Racial underrepresentation in dermatological datasets leads to biased machine learning models and inequitable healthcare. *3*, 42–47.

Dwork, C., Hardt, M., Pitassi, T., Reingold, O., & Zemel, R. (2012). Fairness through awareness. *Proceedings of the 3rd Innovations in Theoretical Computer Science Conference*, 214–226. https://doi.org/10.1145/2090236.2090255
Mehrabi 48.

Farnadi, G., Babaki, B., & Getoor, L. (2018). Fairness in relational domains. *Proceedings of the 2018 AAAI/ACM Conference on AI, Ethics, and Society*, 108–114. https://doi.org/10.1145/3278721.3278733
Mehrabi 50.

Friedman, B., & Nissenbaum, H. (1996). Bias in computer systems. *ACM Trans. Inf. Syst.*, *14*(3), 330–347. https://doi.org/10.1145/230538.230561
Mehrabi 53.

Fry, A., Littlejohns, T. J., Sudlow, C., Doherty, N., Adamska, L., Sprosen, T., Collins, R., & Allen, N. E. (2017). Comparison of sociodemographic and health-related characteristics of UK biobank participants with those of the general population. *American Journal of Epidemiology*, *186*(9), 1026–1034. https://doi.org/10.1093/aje/kwx246
Mehrabi 54.

Gottfrois, P., Gröger, F., Andriambololoniaina, F. H., Amruthalingam, L., Gonzalez-Jimenez, A., Hsu, C., Kessy, A., Lionetti, S., Mavura, D., Ng'ambi, W., Ngongonda, D. F., Pouly, M., Rakotoarisaona, M. F., Rapelanoro Rabenja, F., Traoré, I., & Navarini, A. A. (2024). Passion for dermatology: Bridging the diversity gap with pigmented skin images from sub-saharan africa. *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*, 703–712.

Grgic-Hlača, N., Zafar, M. B., Gummadi, K. P., & Weller, A. (2016). The case for process fairness in learning: Feature selection for fair decision making
Mehrabi 61.

Hajian, S., & Domingo-Ferrer, J. (2013). A methodology for direct and indirect discrimination prevention in data mining [Conference Name: IEEE Transactions on Knowledge and Data Engineering]. *IEEE Transactions on Knowledge and Data Engineering*, *25*(7), 1445–1459. https://doi.org/10.1109/TKDE.2012.72
Mehrabi 62.

Hardt, M., Price, E., Price, E., & Srebro, N. (2016). Equality of opportunity in supervised learning. *Advances in Neural Information Processing Systems*, *29*. Retrieved March 16, 2025, from https://proceedings.neurips.cc/paper/2016/hash/9d2682367c3935defcb1f9e247a97c0d-Abstract.html
Mehrabi 63.

Hargittai, E. (2007). Whose space? differences among users and non-users of social network sites. *Journal of Computer-Mediated Communication*, *13*(1), 276–297. https://doi.org/10.1111/j.1083-6101.2007.00396.x
Mehrabi 64.

HP, S. (2022, November 1). *Sampling — statistical approach in machine learning* [Analytics vidhya]. Retrieved March 28, 2025, from https://medium.com/analytics-vidhya/sampling-statistical-approach-in-machine-learning-4903c40ebf86

Jayawickrama, T. D. (2021, February 1). *Community detection algorithms* [Medium]. Retrieved March 24, 2025, from https://medium.com/data-science/community-detection-algorithms-9bd8951e7dae

Kearns, M., Neel, S., Roth, A., & Wu, Z. S. (2018). Preventing fairness gerry-
  mandering: Auditing and learning for subgroup fairness [ISSN: 2640-3498].
  Proceedings of the 35th International Conference on Machine Learning,
  2564–2572. Retrieved March 16, 2025, from https://proceedings.mlr.press/
  v80/kearns18a.html
  Mehrabi 79.

Kearns, M., Neel, S., Roth, A., & Wu, Z. S. (2019). An empirical study of rich
  subgroup fairness for machine learning. *Proceedings of the Conference on
  Fairness, Accountability, and Transparency*, 100–109. https://doi.org/10.
  1145/3287560.3287592
  Mehrabi 80.

Kusner, M. J., Loftus, J., Russell, C., & Silva, R. (2017). Counterfactual fairness.
  Advances in Neural Information Processing Systems, *30*. Retrieved March
  16, 2025, from https://proceedings.neurips.cc/paper_files/paper/2017/
  hash/a486cd07e4ac3d270571622f4f316ec5-Abstract.html
  Mehrabi 87.

Lerman, K., & Hogg, T. (2014). Leveraging position bias to improve peer recom-
  mendation [Publisher: Public Library of Science]. *PLOS ONE*, *9*(6), e98914.
  https://doi.org/10.1371/journal.pone.0098914
  Mehrabi 93.

Manrai, A. K., Funke, B. H., Rehm, H. L., Olesen, M. S., Maron, B. A., Szolovits,
  P., Margulies, D. M., Loscalzo, J., & Kohane, I. S. (2016). Genetic mis-
  diagnoses and the potential for health disparities. *N Engl J Med*, *375*(7),
  655–665. https://doi.org/10.1056/NEJMsa1507092
  Mehrabi 98.

Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A
  survey on bias and fairness in machine learning [Publisher: ACMPUB27New
  York, NY, USA]. *ACM Computing Surveys (CSUR)*. https://doi.org/10.
  1145/3457607

Mester, T. (2017, August 28). *Statistical bias types explained - part2 (with ex-
  amples)* [Data36]. Retrieved March 22, 2025, from https://data36.com/
  statistical-bias-types-examples-part2/

Mester, T. (2022, May 16). *Statistical bias types explained (with examples) - part1*
  [Data36]. Retrieved March 8, 2025, from https://data36.com/statistical-
  bias-types-explained/

Montoya, L. N., Roberts, J. S., & Hidalgo, B. S. (2025). Towards fairness in AI
  for melanoma detection: Systemic review and recommendations. In K. Arai
  (Ed.), *Advances in information and communication* (pp. 320–341). Springer
  Nature Switzerland. https://doi.org/10.1007/978-3-031-84460-7_21
  2025.

Mustard, D. B. (2003). Reexamining criminal behavior: The importance of omitted
  variable bias. *The Review of Economics and Statistics*, *85*(1), 205–211.
  https://doi.org/10.1162/rest.2003.85.1.205
  Mehrabi 114.

Olteanu, A., Castillo, C., Diaz, F., & Kıcıman, E. (2019). Social data: Biases,
  methodological pitfalls, and ethical boundaries [Publisher: Frontiers]. *Front.*

Big Data, *2*. https://doi.org/10.3389/fdata.2019.00013
Mehrabi 120.

Pala, P., Bergler-Czop, B. S., & Gwiżdż, J. M. (2020). Teledermatology: Idea, benefits and risks of modern age – a systematic review based on melanoma. *Postepy Dermatol Alergol*, *37*(2), 159–167. https://doi.org/10.5114/ada.2020.94834

Riegg, S. K. (2008). Causal inference and omitted variable bias in financial aid research: Assessing solutions [Publisher: Johns Hopkins University Press]. *The Review of Higher Education*, *31*(3), 329–354. Retrieved March 16, 2025, from https://muse.jhu.edu/pub/1/article/232773
Mehrabi 131.

Romani, L., Whitfeld, M. J., Koroivueta, J., Kama, M., Wand, H., Tikoduadua, L., Tuicakau, M., Koroi, A., Ritova, R., Andrews, R., Kaldor, J. M., & Steer, A. C. (2017). The epidemiology of scabies and impetigo in relation to demographic and residential characteristics: Baseline findings from the skin health intervention fiji trial. *Am J Trop Med Hyg*, *97*(3), 845–850. https://doi.org/10.4269/ajtmh.16-0753

Shankar, S., Halpern, Y., Breck, E., Atwood, J., Wilson, J., & Sculley, D. (2017, November 22). No classification without representation: Assessing geodiversity issues in open data sets for the developing world. https://doi.org/10.48550/arXiv.1711.08536
Mehrabi 142Comment: Presented at NIPS 2017 Workshop on Machine Learning for the Developing World.

Suresh, H., & Guttag, J. (2021). A framework for understanding sources of harm throughout the machine learning life cycle. *Proceedings of the 1st ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization*, 1–9. https://doi.org/10.1145/3465416.3483305
Mehrabi 144.

Taylor, C. (2023, April 5). *Unbiased and biased estimators* [ThoughtCo] [Section: ThoughtCo]. Retrieved April 5, 2025, from https://www.thoughtco.com/what-is-an-unbiased-estimator-3126502

Verma, S., & Rubin, J. (2018). Fairness definitions explained. *Proceedings of the International Workshop on Software Fairness*, 1–7. https://doi.org/10.1145/3194770.3194776
Mehrabi 149.

Vickers, S. M., & Fouad, M. N. (2014). An overview of EMPaCT and fundamental issues affecting minority participation in cancer clinical trials. *Cancer*, *120*(0), 1087–1090. https://doi.org/10.1002/cncr.28569
Mehrabi 150.

Wang, T., & Wang, D. (2014). Why amazon's ratings might mislead you: The story of herding effects [Publisher: Mary Ann Liebert, Inc., publishers]. *Big Data*, *2*(4), 196–204. https://doi.org/10.1089/big.2014.0063
Mehrabi 151.

Young, A. T., Xiong, M., Pfau, J., Keiser, M. J., & Wei, M. L. (2020). Artificial intelligence in dermatology: A primer. *Journal of Investigative Dermatology*,

140(8), 1504–1512. https://doi.org/10.1016/j.jid.2020.02.026
209 citations.

Zhao, J., Wang, T., Yatskar, M., Ordonez, V., & Chang, K.-W. (2017, July 29).
Men also like shopping: Reducing gender bias amplification using corpus-
level constraints. https://doi.org/10.48550/arXiv.1707.09457
Mehrabi 167 Comment: 11 pages, published in EMNLP 2017.

Zhao, J., Wang, T., Yatskar, M., Ordonez, V., & Chang, K.-W. (2018, April 18).
Gender bias in coreference resolution: Evaluation and debiasing methods.
https://doi.org/10.48550/arXiv.1804.06876
Mehrabi 168, Comment: NAACL '18 Camera Ready.

Projektspezifisch können weitere Dokumentationsteile angefügt werden wie: Aufgabenstellung, Projektmanagement-Plan/Bericht, Testplan/Testbericht, Bedienungsanleitungen, Details zu Umfragen, detaillierte Anforderungslisten, Referenzen auf projektspezifische Daten in externen Entwicklungs- und Datenverwaltungstools etc.

TODO: check the gls all unused.